

# Execution Report #5

The register runs dry

B5 (engine substitution), B7 (escalation band), B8 (specialization boundary),  
B9 (context paging), B4 (deontic fragment), and the consistency-radius theorem —  
the entire blocking-before-arXiv register, discharged

Prepared for Erich Owens

August 2026

**Express lane.** Six items, six scripts, all green at seed 20260816: the reputation market unravels without engine attestation and flips to the planner's rule with it (§1); the escalation debit has a unique separating threshold and a two-walled tuning band that can be honestly empty (§2); the whitepaper's specialization threshold is *falsified* and replaced by the exact Erlang-C boundary, with the succession rule priced (§3); context paging degrades only linearly under a corrupted pin oracle — if and only if pins are re-derived on touch (§4); commitment-conflict detection is polynomial inside the Horn fragment and NP-complete one step outside (§5); and the sheaf detector's residual is now a *certified lower bound* on the injected lie, with localization and complexity theorems (§6). With these, the proof-debt register's blocking-before-arXiv list (A3, A4, B5, B7, B8, B9) is clear. Every number is tagged [verified] (externally recomputable) or [internal] (regenerates from the named script, seed 20260816).

## 1 B5 — The engine behind the reputation

**The story.** An identity builds a glowing record running a strong engine, then quietly swaps a cheap one in behind the same name. Buyers see the reputation; they do not see the engine. This is the resume-button feature's named attack — and it is Akerlof's used-car market, run *inside one identity*.

*The idea in one breath: unattested, the price can never depend on the actual engine, so swapping always pays and pooling on quality dies; attest the engine on every witnessed outcome and the substitution incentive flips to the planner's own rule — and the same three checks make resurrection across providers sanction-sound.*

**Theorem (unraveling).** With engines  $\theta_H > \theta_L$  at costs  $c_H > c_L$  and the engine unobserved, the swap gain  $(p - c_L) - (p - c_H) = \Delta c > 0$  is price-independent: no belief supports pooling on  $\theta_H$  and nothing exists to separate on. Committed  $\theta_H$  sellers participate iff  $\mu \geq \mu^* = (c_H - \theta_L)/\Delta\theta$ ; below it the death spiral rests at  $\theta_L$ . Deterrence without attestation needs an audited stake  $qW \geq \Delta c$ .

**Theorem (IC flip).** With daemon-attested engine ids, reputation keys on (principal, engine) and the swap gain becomes  $\Delta c - \Delta\theta$ : substitution pays iff the weak engine is socially better, at zero audit stake.

**Theorem (resurrection soundness).** Sanction-respecting reputation (Def. III.6.1) survives provider migration under (i) credential+continuity-witness verification, (ii) engine attestation in successor outcomes, (iii) open commitments closed or renegotiated with escrow delta — and each dropped clause admits a shortest-trace escape: whitewash-by-

resurrection, engine-history shedding, liability stranding.

**Numbers.**  $\theta_H=1, \theta_L=0.4, c_H=0.5, c_L=0.2$ : swap gain 0.3 at every price,  $\mu^* = 1/6$ , attested keep 0.5 vs swap 0.2 [verified]. Sweeps: 4000 instances  $\times$  8 prices, 0 pooling survivals; thresholds match simulation exactly; 2000 heterogeneous spirals monotone to the Akerlof fixed point ( $\mu_0=0.5 \rightarrow 1/3$ ;  $\mu_0=0.2 \rightarrow$  total unraveling); migration machine: 747 states to depth 7, 0 Def. III.6.1 violations intact; mutants caught in 1, 2, 2, 2 steps [internal, `b5_engine_substitution.py`].

**The wrong turn, reported.** A first pass keyed *sanctions* like scores, per (identity, engine) — and the intact protocol failed Def. III.6.1 in one step, no migration needed: the sibling engine key is a free escape hatch. The fix is the paper’s one-liner: price per (principal, engine), sanction per principal — one rule closing engine substitution and skill-versioning Gap #4 together.

**Honest boundary.** Two periods, no continuation value beyond the stake; attestation soundness is a deployment (TEE) assumption, per the declined-claims list; the migration result is bounded small-model checking of the protocol spec, which the daemon must refine; Def. III.6.1 is score-only — escrow enters as attribution, not conservation.

## 2 B7 — The costly-escalation debit: threshold equilibrium and its tuning band

**The story.** An agent privately observes urgency  $u \sim F$  on  $[0, 1]$ ; escalating surfaces the item on the operator’s reserved distress lane, where it is validated with probability  $V(u)$  and yields benefit  $b(u)$ ; a *dismissed* escalation debits standing by  $\delta$ , worth  $w$  per unit of surfacing weight. Crying wolf has a price — but what price is right?

*Monotone benefit and validation make the escalation payoff single-crossing, so a unique threshold type separates escalators from the silent; and between the alarm-fatigue wall and the distress-silencing wall the feasible debit is an interval — possibly empty, in which case no tuning saves an under-provisioned channel.*

**Theorem (threshold equilibrium; single crossing).** With  $\Pi_\delta(u) = b(u) - \delta w(1 - V(u))$ , monotone  $b$  and  $V$  make  $\Pi_\delta$  strictly increasing, so with  $\Pi_\delta(0) < 0 < \Pi_\delta(1)$  there is a unique  $u^*(\delta)$  solving  $b(u^*) = \delta w(1 - V(u^*))$ , and the unique best response is “escalate iff  $u \geq u^*(\delta)$ ” — a separating equilibrium in the Spence mold.  $u^*$  is nondecreasing in  $\delta$  with interior slope  $w(1 - V(u^*)) / (b'(u^*) + \delta w V'(u^*))$ . For  $u \sim \text{Unif}[0, 1]$ ,  $b(u) = u$ ,  $V(u) = u$ :  $u^*(\delta) = \delta w / (1 + \delta w)$ .

**Theorem (the tuning band).** With alarm load  $L(\delta) = \int_{u^*(\delta)}^1 [c_{\text{att}} + C_{\text{fa}}(1 - V(u))] dF$  nonincreasing and miss loss  $\text{ML}(\delta) = C_{\text{miss}}(F(u^*(\delta)) - F(u_{\text{crit}}))^+$  nondecreasing, the feasible set  $\{\delta : L \leq A, \text{ML} \leq M\}$  is an interval  $[\delta_{\min}, \delta_{\max}]$ , nonempty iff  $L(\delta_{\max}) \leq A$ ; here  $u_{\text{crit}} = (c_{\text{att}} + C_{\text{fa}}) / (C_{\text{miss}} + C_{\text{fa}})$  is exactly R3’s inspect threshold.

**Numbers.** Closed forms match bisection ( $u^* = \delta w / (1 + \delta w)$ ; Lambert- $W$  for the exponential family) [verified]; single crossing and monotone  $u^*(\delta)$ : 0/4000 violations over random monotone families; the non-monotone attack breaks the threshold in 1172/2000 draws (witness:  $u = 0.681$  escalates while  $u = 0.751$  is silent) [internal, `b7_escalation_band.py`]. R3’s constants ( $C_{\text{miss}}=100$ ,  $c_{\text{att}}=1$ ,  $C_{\text{fa}}=5$ ),  $w=1$ ,  $M=1$ : budget  $A = 3.3$  gives  $\delta \in [0.035, 0.072]$ ; budget  $A = 2.0$  gives  $\delta_{\min} = 0.396 > \delta_{\max} = 0.072$  — *empty*. The zero-debit mutant blows the fatigue budget ( $L(0) = 3.5$ , the Cvach habituation regime);  $\delta = 1$  silences 0.44 of genuine-distress mass (miss loss 44.3) yet passes the fatigue wall alone — one-sided tuning would miss it [internal].

**Honest boundary.** Monotone  $V$  is the condition earning its keep: non-monotone validation produces multi-interval escalation sets — measure the validation curve before trusting any  $\delta$ . Empty-band regimes are real; there the fix is attention capacity or evidence quality, not tuning. Static one-shot model: no reputation dynamics of  $w$ , no operator strategic dismissal, no flood externality (R7/A2’s job).

### 3 B8 — The specialization boundary, and the price of the succession rule

**The story.** Should an invariant-critical function have one deeply skilled owner or a pool of generalists? The whitepaper proposed a threshold. The sweep ran first — and the proposed threshold is *false in both directions*: it certifies specialists who lose at low utilization and rejects specialists who win at high utilization.

*The exact boundary is Erlang-C: sole ownership pays iff the skill premium clears  $g(\rho, c) = c\rho + c(1 - \rho)/(c(1 - \rho) + C(c, \rho))$  — which saturates at  $c$ , never diverges; and once the owner’s death-rate-times-succession-time product exceeds  $D^*$ , no skill premium rescues sole ownership at all.*

**Theorem (queueing specialization boundary).** Requests arrive Poisson( $\lambda$ ); a sole specialist serves as  $M/M/1$  at rate  $\mu_s$ , the pool as  $M/M/c$  at  $\mu_g$  each, utilization  $\rho = \lambda/(c\mu_g)$ ; accountability value  $A$ , waiting cost  $w$ . With  $C(c, \rho)$  the Erlang-C delay probability, sole ownership weakly dominates iff

$$\frac{\mu_s}{\mu_g} \geq g_A(\rho, c) = c\rho + \frac{1}{1 + \frac{C(c, \rho)}{c(1 - \rho)} + \frac{A\mu_g}{w\lambda}},$$

reducing at  $A = 0$  to  $g(\rho, c) = c\rho + \frac{c(1 - \rho)}{c(1 - \rho) + C(c, \rho)}$ , with  $g(\rho, 1) = 1$ ,  $g \rightarrow c$  as  $\rho \rightarrow 1$ , and  $g(\rho, 2) = 1 + 2\rho - \rho^2$  in closed form.

**Theorem (price of the succession rule).** Specialist dies at rate  $\xi$ , succeeded after  $\text{Exp}(\eta)$  spin-up, work resuming, requests queueing throughout. With  $\mu_{\text{eff}} = \mu_s\eta/(\xi + \eta)$ :  $W = (1 + \mu_s\xi/(\xi + \eta)^2)/(\mu_{\text{eff}} - \lambda)$  exactly, with infimum  $W_\infty = \xi/(\eta(\xi + \eta))$  as  $\mu_s \rightarrow \infty$ : no skill buys back residual downtime. Pooling dominates at *every* skill premium iff  $\xi/\eta > D^* = \eta K/(1 - \eta K)$ ,  $K = A/(w\lambda) + W_{\text{pool}}$ .

**Numbers.** The proposed threshold  $\tilde{g} = 1 + (c-1)\rho/(1-\rho)$  is falsified both directions:  $c=2, \rho=0.1$  false-certify (sim  $z = -15.7$ );  $c=2, \rho=0.5$  false-reject ( $z = +32$ ); crossing at  $\rho = (3 - \sqrt{5})/2 \approx 0.382$ ;  $\tilde{g}$  diverges where  $g$  saturates ( $\rho=0.9, c=8$ :  $g = 7.73$  vs  $\tilde{g} = 64$ ) [internal, `b8_specialization.py`].  $W$  closed form vs matrix-geometric  $2 \times 10^{-13}$ , vs truncated CTMC  $10^{-14}$ , vs Gillespie  $|z| \leq 0.8$ ; 60-instance sweep, 0 sign violations; succession instance  $\eta = 0.25$ :  $D^* = 0.350$ , and at  $\xi/\eta = 2$  even  $\mu_s = 10^8$  loses [internal]. The breakdown-forgetting mutant is caught by the high- $\xi$  canary. A common misread to preempt:  $D^*$  does not say specialists are bad — it prices exactly when a succession plan (large  $\eta$ ) keeps sole ownership viable.

**The wrong turn, reported.** The naive decomposition  $W = 1/(\mu_{\text{eff}} - \lambda) + W_\infty$  is refuted (4.17 vs true 8.17 at the canary instance): downtime also builds queue, amplifying the residual by  $\mu_{\text{eff}}/(\mu_{\text{eff}} - \lambda)$ .

**Honest boundary.** Mean-cost comparison only (no tail/SLA percentiles); FCFS, exponential everything; breakdowns-at-all-times with preemptive resume (the death/heartbeat model);  $A$  and  $w$  are policy prices the theorem constrains, not picks.

## 4 B9 — Context paging with an untrusted pin oracle

**The story.** The resident context is a cache of spans; a feature oracle marks some spans load-bearing (pinned). The oracle is sometimes wrong — at rate  $\varphi$ , the same forgeable-feature budget R1 already carries. How badly does paging degrade?

*Linearly, and only linearly: with Young’s Landlord on the honestly-unpinned region and pins re-derived on every fetch, a corrupted consultation costs at most one refetch — but trust a pin blindly and a single corrupted bit costs  $\Theta(N)$ .*

**Theorem (import; Young 2002).** Landlord with cache  $k$  over spans of size  $s(f)$  and refetch cost  $c(f)$  satisfies  $\text{cost}_{LL}(\sigma) \leq \frac{k}{k-h+1} \text{OPT}_h(\sigma)$  for every  $h \leq k$  — the weighted file-caching transfer variable-size context spans need.

**Theorem (linear  $\varphi$ -degradation).** Unit cost density  $c(f) = s(f)$  (the context-window regime); honest pinned set  $P$ ; oracle corrupted on  $C$  consultations ( $\mathbb{E}[C] = \varphi N$ , obviously), pins re-derived on every fetch (*repair-on-touch*). Then pin-augmented Landlord pays, for every  $p \leq h \leq k$ ,

$$\text{Cost} \leq c(P) + \frac{k-p}{k-h+1} \text{OPT}_{h-p}(\sigma_{\text{unpinned}}) + \varphi N c_{\text{refetch}},$$

an *additive*, linear  $\varphi$  term: each corruption perturbs one eviction pass and is charged at most one refetch (missed pin: repair-on-touch buys it back; mis-pinned junk: the displaced volume refetches for at most  $s(g) \leq c_{\text{refetch}}$ ).

**Numbers.** Import vs exact OPT (subset-DP): 122 (instance,  $h$ ) pairs, 0 violations, cyclic thrash instances at exact equality [internal, `b9_context_paging.py`]. Degradation: 50 instances at  $\varphi \in \{0.1, \dots, 0.4\}$ , all  $h$ , 0 violations; cost-vs- $\varphi$  slope 22.8 vs ceiling  $N c_{\text{refetch}} = 1200$  (1.9%;  $R^2 = 0.994$ ); per-corruption charge  $\leq 0.100 c_{\text{refetch}}$  over 120 trials. An *adaptive* corruptor that sees cache state maxed at  $0.778 c_{\text{refetch}}$  per corruption — inside the bound; linearity did not break. Mutation: blind pin-trust turns *one* corruption into 317 refetch-equivalents at  $N=1920$  vs the graceful policy’s exactly one refetch [internal].

**The wrong turn, reported.** A corruption can *help* (extra cost  $< 0$ ) when the falsely-unpinned span was crowding a thrashing working set — the bound is one-sided by design; the final mutation scenario isolates the pin effect.

**Honest boundary.** Unit cost density is load-bearing for the mis-pin charge. Repair-on-touch is *the* load-bearing guard — without it damage is unbounded (staleness is priced charitably at refetch cost; acting on absent load-bearing context is a correctness failure). The oblivious hypothesis was not the boundary for the additive term (greedy-adaptive stayed inside; the oblivious/adaptive gap lives in the multiplicative  $O(\log k)$  of randomized marking); a stronger adaptive adversary is untested. The sweep is evidence at small scale; the second-order credit step rests on the Landlord term’s slack, verified per instance.

## 5 B4 — The tractable conflict fragment and its frontier

**The story.** The condominium harbor’s lookout must answer, at proposal time, whether an incoming commitment conflicts with the accepted set. Cheap detection is a language-design choice: the commitment language buys it by construction.

*Inside Horn rules, ground intervals, and difference constraints, conflict detection is a polynomial, witness-producing decision; allow one disjunctive obligation and it is NP-complete — the fragment boundary is not taste but the exact price of proposal-time checking.*

**Theorem (membership).** In the fragment  $\mathcal{L}_c$  (facts; definite Horn rules with  $\perp$  heads allowed; deontic rules body  $\rightarrow M(\text{act}, s, [t_1, t_2])$ ,  $M \in \{O, F\}$ ; exclusive interval claims; difference constraints  $x_j - x_i \leq c$ ), conflict detection is decidable in  $O(H + T \log T + C \log C + V \cdot E)$  with a polynomial witness per conflict: Dowling–Gallier unit propagation computes the least Horn model (= statuses forced in every model, by monotonicity); a keyed sweep-line finds every O/F and claim overlap; Bellman–Ford returns a negative cycle iff the deadline system is infeasible.

**Theorem (frontier).** Extend  $\mathcal{L}_c$  with disjunctive obligations  $O(l_1 \vee \dots \vee l_k)$  under discharge-choice semantics: conflict-freedom is NP-complete (membership via the selection certificate; hardness from 3-SAT — clause  $\mapsto$  disjunctive obligation over selector literals, variable  $\mapsto$  an O/F pair on identical scope and interval).

**Numbers.** 3000 random in-fragment policy sets vs a brute-force oracle (all  $2^5$  Horn models intersected; exhaustive integer search over the complete potential box): 885 conflicts (121  $\perp$ , 195 O/F, 484 claims, 213 temporal), 149 reachable only through Horn propagation, **0 disagreements**; least model = intersection-of-all-models on all 3000 [internal, `b4.deontic.fragment.py`]. The reduction verified both directions on 8 SAT + 8 UNSAT instances, 16/16. Mutation: removing Horn propagation misses the four-line indirect witness (deploy  $\rightarrow$  prod access  $\rightarrow$  obliged write vs standing prohibition) and 100/885 sweep conflicts — propagation is load-bearing at scale.

**Honest boundary.** Completeness is relative to least-model semantics (monotone bodies; negation-as-failure leaves the fragment). Ground atoms and intervals only — unification and symbolic endpoints coupled to difference variables are outside what is proven. The batch bound is proven; the incremental form is a constants improvement. The checker certifies the norm *system*, not the norms: a bad policy consistently held is consistent.

## 6 The consistency-radius theorem (W8 tail)

**The story.** The R6 harness validated the completion residual  $r$  statistically. What was missing was the theorem: what exactly does  $r > 0$  prove, how large a lie does it certify, where does it point, and what does it cost to compute?

*$r$  is the exact minimum edge-space perturbation any adversary must have injected, it localizes to cycles through the equivocator, and it computes as at most  $L$  scalar graph-Laplacian solves — with the single-equivocator condition now proved to be the honest scope, because coalitions on a cycle can cancel to zero.*

**Theorem CR-1 (soundness).**  $r > 0$  iff no global section explains the compared+relayed data; honest executions give  $r = 0$  exactly, for every visibility pattern. Every offset pattern consistent with the data satisfies  $\|\varepsilon_K\| \geq r$ , with equality at the completion. Single equivocator  $q$ :  $\sigma_{\min}^+(M_q) \text{dist}(o, \ker M_q) \leq r \leq \text{dist}(o, \ker M_q)$ , and for a single-edge lie of size  $s$  in coordinate  $c$ ,  $r = |s| \sqrt{1 - R_{\text{eff}}^{G_c}(e)}$  — on  $C_n$ ,  $|s|/\sqrt{n}$ : the harness’s 1.2247 is  $3\sqrt{1 - 5/6}$ .

**Theorem CR-2 (localization).** Noise-free, single equivocator: the residual’s support lies on cycles of the coordinate subgraphs through  $q$  (electrical-flow support argument) — the harness’s 200/200 is now a theorem; with noise, guaranteed while  $2\|\eta_K\| < \max_f \|(\Pi_K \varepsilon_K)_f\|$ .

**Theorem CR-3 (complexity).**  $r$  is one sparse least-squares that decomposes into at most  $L$  scalar graph-Laplacian systems:  $\tilde{O}(|E|L \log(1/\epsilon))$  with SDD solvers. (Bach’s #P-hardness concerns coherent sheaves on projective space, not finite cellular sheaves over

$\mathbb{R}$ .)

**Numbers.** Soundness 0/600 violations, achievability 600/600; the effective-resistance closed form exact on  $C_4$ – $C_{12}$  and 200 random graphs; the  $\sigma_{\min}^+$  lower bound tight under the adversarial offset (0.4082); a uniform (kernel) lie of norm 5 gives  $r = 2 \times 10^{-14}$  — invisible in principle, correctly not flagged [internal, `sheaf_consistency_radius.py`]. Localization 200/200 (two-cycles topology) + 128/128 (severed expanders); the noise boundary is measured:  $\geq 95\%$  up to  $\sigma_\eta/\sigma_{\text{eq}} \approx 0.1$ , 76.5% at 0.2. Timing slope vs  $|E|$ : CG 0.65 (linear-ish), dense 1.88. The sweep found the promised crack: two coordinated liars on a common cycle cancel to  $r = 6 \times 10^{-15}$  while each alone is detectable at  $|s|/\sqrt{8}$  — the detection bound is single-equivocator; soundness ( $r$  never overstates) is unconditional. The D2-reintroduction mutant fires on a severed lie exactly as CR-1(i) predicts.

**Honest boundary.** Detection and localization are single-equivocator: coalition-coboundary lies are consistent counterfactual worlds, provably dark.  $r$  attributes nothing (a  $+o$  lie by  $q$  toward  $w$  and a  $-o$  lie by  $w$  toward  $q$  yield identical observations — signatures attribute, cohomology localizes). Vanishing  $r$  is not an all-clear (uniform lies, severed edges, the Carù abelianization gap). Under noise, soundness needs a threshold above the honest floor.

## 7 Program status

| Item                             | Status      | Headline                                        | Feeds         |
|----------------------------------|-------------|-------------------------------------------------|---------------|
| A1–A4, A6, A7, B1–B3, B6, C0, C1 | done        | Execution Reports 1–4                           | Papers 1–4, 8 |
| Zoom theorem; sheaf harness      | done        | Paper 1 closed; verdict COMMIT                  | Papers 1, 7   |
| <b>B5</b>                        | <b>done</b> | unraveling; IC flip; resurrection soundness     | Paper 5       |
| <b>B7</b>                        | <b>done</b> | threshold equilibrium; tuning band              | Paper 1 orbit |
| <b>B8</b>                        | <b>done</b> | exact Erlang-C boundary; $D^*$ succession price | Paper 6       |
| <b>B9</b>                        | <b>done</b> | Landlord import; linear $\varphi$ -degradation  | Paper 1 orbit |
| <b>B4</b>                        | <b>done</b> | polynomial fragment; NP-complete frontier       | Paper 6       |
| <b>CR-1/2/3</b>                  | <b>done</b> | $r$ certifies, localizes, computes              | Paper 7       |

The blocking-before-arXiv register (A3, A4, B5, B7, B8, B9) is **clear**. Remaining: Papers 5–7 assemblies (all inputs now in hand), C2 (minimum-viable take-rate), the lifts ( $C0 \rightarrow \text{Apalache}$ ,  $C1 \rightarrow \text{Isabelle}$ ), and the whitepaper folds (thm:specialization must adopt the exact  $g$ ; the context-paging and escalation status notes upgrade to executed).

*Now you try.* Compute  $D^*$  for a sole-owned role in your own deployment: measure its request rate  $\lambda$ , the pool’s service rate  $\mu_g$ , the role’s historical death rate  $\xi$  (departures, deprecations), and your successor spin-up time  $1/\eta$ ; if  $\xi/\eta$  exceeds  $\eta K/(1 - \eta K)$ , no skill premium rescues sole ownership — write the succession plan first.

*Reproducibility.* Six self-contained scripts (`b5_engine_substitution.py`, `b7_escalation_band.py`, `b8_specialization.py`, `b9_context_paging.py`, `b4_deontic_fragment.py`, `sheaf_consistency_radius.py`); every [internal] number regenerates deterministically at seed 20260816, every mutation prints its witness, and each script exits nonzero on any failed check. All six run in the harbor-results-estate CI job.