

Regimented or Enforced

The Controllability Boundary for Agent Governance

Paper 2 of the Harbor program: a supervisory-control characterization
of what an agent runtime can prevent

Prepared for Erich Owens

August 2026

Abstract

An agent hypervisor — a runtime that mediates an LLM agent’s file writes, network egress, tool executions, pushes, and spawns — is asked to *enforce* governance policies. Some it can enforce by prevention: the forbidden thing never happens. Others it can only enforce by detection and compensation: the thing happens, and the runtime notices afterward. We show the boundary between the two is not an engineering contingency but a theorem: a prefix-closed safety policy is *regimentable* (preventable pre-effect) if and only if it is controllable in the sense of Ramadge–Wonham supervisory control theory, with respect to the uncontrollable event set of an LLM (model-internal steps and in-context reads). Schneider’s classical result — execution monitors enforce exactly safety properties — is the ambient case of this characterization; controllability is the refinement forced by an alphabet the monitor only partly mediates. The characterization is executable: a product-automaton checker classifies a realistic policy set (egress, push, write, exec, spawn: regimentable; token emission, context reads, internal plans, confident falsehoods: detect-only forever), and settles the decisive compound case — “no egress after reading a secret” *is* regimentable, because the policy permits the uncontrollable trigger and gates only the controllable effect. That single classification is the design theorem behind clean-room agent computation: gate the channel, never the token.

Express lane. *A governance policy can be enforced by prevention iff no legal prefix of agent behavior can be pushed out of the policy by an event the runtime cannot intercept — Ramadge–Wonham controllability with respect to $\Sigma_u = \{\text{model-internal steps, in-context reads}\}$; formal statement and proof in the box of §4, machine-checked policy table in §5. Experts: read the box, the table, and §6; the rest is scaffolding.*

1 The scene: two kinds of “the runtime enforces it”

A team runs coding agents behind a single-writer daemon. Every file write, network call, subprocess spawn, and `git push` passes through the daemon’s effect boundary, where policy is checked before the effect lands. The governance document lists rules: *no force-push to main*; *no network egress from sandboxed jobs*; *no writes outside the granted scope*; and also *no confident falsehoods in status reports*; *never incorporate the customer’s secret into your plan*.

The operations lead asks the obvious question: which of these does the runtime *guarantee*, and which does it merely *watch for*? The first three feel different from the last two, and every practitioner senses it. The force-push rule is airtight — the daemon owns the git remote credentials, and a push that never happens needs no forgiveness. The confident-lie rule is not airtight and never will be: the lie is composed inside the model’s forward pass, and by the time any mediator sees tokens, the lie already exists. The practitioner’s instinct is correct. This paper proves it is not an instinct but an instance of a thirty-nine-year-old theorem, and extracts the exact boundary.

The idea in one breath.

Split the agent’s event alphabet into what the runtime can refuse (mediated effects) and what it cannot (model-internal steps); a safety policy is preventable iff no uncontrollable event can carry a legal history out of the policy — and that is precisely Ramadge–Wonham controllability.

Intuition: the bouncer. A club has one door and a bouncer. The bouncer can refuse *entry* — entry is controllable, mediated by the door. The bouncer cannot refuse a patron’s *thought* — thought is uncontrollable, unmediated by anything the club owns. A door policy (“no entry after 2am”) is enforceable by prevention. A thought policy (“no ill intent inside”) is enforceable only by observation and ejection — detection and compensation. Now the structural mapping (Figure 1): the daemon’s effect boundary is the door; Σ_c (fs_write, net_egress, exec_tool, git_push, spawn_child) is entry; Σ_u (model_emit_token, in_context_read, internal_plan) is thought. The analogy maps relations, not surface features, so it earns a prediction: a *compound* rule mentioning thought but gating only the door — “no entry for anyone who arrived after the fight started” — should still be enforceable by prevention, because the bouncer needs only to *observe* the trigger and *refuse* the door. Section 6 confirms exactly this for agents.

The predictable misread, preempted: the analogy does *not* say internal events are invisible — tokens are eminently loggable after emission. Uncontrollable means *unrefusable before the fact*, not unobservable after it. (When events are also unobservable, a second, stricter theory applies — §8.)

2 The model

Events and the mediation boundary. Fix a finite alphabet $\Sigma = \Sigma_c \uplus \Sigma_u$ of agent events. Σ_c (*controllable*) contains the events that pass through the runtime’s mediation boundary before taking effect, so the runtime can disable any of them at any instant: in our reference alphabet, $\Sigma_c = \{\text{fs_write}, \text{net_egress}, \text{exec_tool}, \text{git_push}, \text{spawn_child}, \text{api_call}\}$. Σ_u (*uncontrollable*) contains the events that occur before any mediator can intervene: $\Sigma_u = \{\text{model_emit_token}, \text{in_context_read}, \text{internal_plan}, \text{hidden_activation}\}$. The split is a fact about the architecture, not a modeling choice: the daemon owns the syscall-and-credential surface and does not own the forward pass.

Plant and policy. The *plant* $L \subseteq \Sigma^*$ is the prefix-closed language of physically possible behaviors; for a general agent we take the universal plant $L = \Sigma^*$ (any interleaving can occur). A *policy* is a prefix-closed language $K \subseteq L$: the set of allowed histories. Prefix closure makes K a *safety* property in the classical sense — a violation is witnessed by a finite prefix and can never be un-violated. $\overline{K}$ denotes prefix closure (so $\overline{K} = K$ here; we keep the bar to match the standard statement).

What “regiment” means. A *regimentation mechanism* is a supervisor $f : \overline{L} \rightarrow 2^{\Sigma_c}$ mapping each history to a set of *disabled controllable* events; the supervised plant executes any event not disabled. Crucially the supervisor’s type says everything: it can disable only members of Σ_c . The mechanism *regiments* K if the closed-loop behavior equals $\overline{K}$ — every allowed history remains reachable and no violating history is reachable. Anything short of that — allowing the violation and then flagging, rolling back, slashing a bond — is *detect-and-compensate*, a legitimate but categorically weaker enforcement mode.

R5 — gate the door, never the thought

hypervisor enforceability = Ramadge–Wonham supervisory control




Figure 1: The relation map. Base domain: a club’s single door — the bouncer refuses entry (controllable) but cannot refuse thought (uncontrollable). Target domain: the daemon’s mediation boundary — the runtime can forbid `fs.write`, `net.egress`, `exec.tool`, `git.push`, `spawn.child`, and cannot forbid token emission, in-context reads, or internal plans. The correspondence arrows carry the theorem: refusable-at-the-boundary $\Leftrightarrow$ preventable (regimentable); internal $\Leftrightarrow$ detect-only forever. The criterion under the map, $\overline{K}\Sigma_u \cap \overline{L} \subseteq \overline{K}$, is Ramadge–Wonham controllability [verified framework].

3 The ambient result, and why the alphabet split refines it

Schneider’s classical theorem says an execution monitor (EM) — a mechanism that watches a single execution and can terminate it — enforces exactly the safety properties [4]. That result is the *ambient* case of ours: it implicitly assumes the monitor mediates *every* event, i.e. $\Sigma_u = \emptyset$. Under complete mediation, controllability is vacuous (there is no uncontrollable event to carry a history out of K), so “safety” is the entire boundary, and Schneider’s theorem is exactly what our characterization degenerates to.

An LLM agent runtime does not enjoy complete mediation. The model’s forward pass emits tokens, reads its context, and forms plans on hardware the daemon schedules but does not gate step-by-step; these events happen and *then* become visible, which in supervisory-control terms makes them uncontrollable. Once $\Sigma_u \neq \emptyset$, safety remains *necessary* for preventability (a liveness violation has no finite witness to block) but is no longer *sufficient*: “never `internal_plan`” is as prefix-closed a safety property as “never `git_push`,” yet only one of them is preventable. The refinement that separates them is Ramadge–Wonham controllability [1, 2], developed for exactly this situation — a plant in which some events (machine breakdowns, arrivals) cannot be disabled by any supervisor. Anderson’s reference-monitor tradition [5] contributes the design obligation (complete mediation *of the controllable alphabet*: every Σ_c event actually passes the boundary), and Rushby’s separation-kernel analysis [6] the standard for verifying that the obligation holds; neither supplies the boundary itself. The boundary is controllability.

4 The theorem

Theorem (Regimentation = controllability). Let $L \subseteq \Sigma^*$ be a prefix-closed plant over $\Sigma = \Sigma_c \uplus \Sigma_u$, and let $K \subseteq L$ be a nonempty prefix-closed safety policy. There exists a regimentation mechanism for K — a supervisor disabling only controllable events whose closed loop realizes exactly $\bar{K}$ — if and only if K is *controllable* with respect to L and Σ_u :

$$\bar{K}\Sigma_u \cap \bar{L} \subseteq \bar{K},$$

i.e. no uncontrollable event, physically enabled after a legal history, exits the policy. When the condition fails, no runtime confined to disabling Σ_c can prevent the violation; the best achievable enforcement is detect-and-compensate, and the largest preventable weakening of the policy is the supremal controllable sublanguage $\sup\mathcal{C}(K)$, which exists and is computable for regular K and L [verified framework, Ramadge–Wonham 1987].

Proof. ($\Leftarrow$) Suppose K is controllable. Define the supervisor $f(s) = \{a \in \Sigma_c : sa \in \bar{L} \setminus \bar{K}\}$ for $s \in \bar{K}$: after any legal history, disable exactly the controllable events that would exit the policy. We show the closed-loop language equals $\bar{K}$, by induction on history length. The empty history lies in $\bar{K}$ (nonempty and prefix-closed). Assume the current history $s \in \bar{K}$. Any event a the closed loop can execute next is enabled in the plant ($sa \in \bar{L}$) and not disabled. If $a \in \Sigma_c$, then $a \notin f(s)$ means $sa \in \bar{K}$ by construction. If $a \in \Sigma_u$, then $sa \in \bar{K}\Sigma_u \cap \bar{L} \subseteq \bar{K}$ by controllability. Either way the closed loop stays in $\bar{K}$, so its language is contained in $\bar{K}$. Conversely every $sa \in \bar{K}$ is reachable: $a \in \Sigma_u$ is never disabled, and $a \in \Sigma_c$ with $sa \in \bar{K}$ is not in $f(s)$; so containment is equality, and f regiments K .

($\Rightarrow$) Suppose K is not controllable: there exist $s \in \bar{K}$ and $\sigma \in \Sigma_u$ with $s\sigma \in \bar{L}$ but $s\sigma \notin \bar{K}$. Let M be any regimentation mechanism for K . Since M realizes exactly $\bar{K}$, the legal history s is reachable under M — M cannot forbid it without cutting $\bar{K}$ down. At s , the plant enables σ , and M cannot disable it: $\sigma \notin \Sigma_c$, and disabling it is outside M ’s type — physically, the event completes before the mediation boundary sees it. So the closed loop under M reaches $s\sigma \notin \bar{K}$, contradicting the assumption that M ’s closed loop is contained in

$\bar{K}$. Hence no regimentation mechanism exists, and any enforcement of K must tolerate the reachable violation $s\sigma$ — observe it after the fact and respond, which is detect-and-compensate by definition. The existence, uniqueness, and computability of $\sup\mathcal{C}(K)$ are imported from Ramadge–Wonham [1]: controllable sublanguages of K are closed under arbitrary union, so a supremal element exists, and for regular languages it is computed by a fixpoint iteration on the product automaton. $\square$

Two remarks on what the proof used. First, the forward direction is constructive and *maximally permissive*: the supervisor disables nothing except the controllable exits, so regimentation costs no legal behavior — the governance rule and the agent’s freedom are not in tension when the rule is controllable. Second, the reverse direction used nothing about the mechanism except its type ($f : \bar{L} \rightarrow 2^{\Sigma_c}$). No cleverness inside the runtime — better prompts, faster monitors, a smarter policy engine — escapes the theorem, because the obstruction is that σ finishes before any decision point the runtime owns.

5 The characterization, made runnable

The controllability test is decidable by a product-automaton search: walk the synchronous product of plant and policy automaton; flag any reachable product state where the plant enables an uncontrollable event that the policy disables. We implemented the test over the reference alphabet and ran it on the governance policy set. The checker¹ produces:

Policy	Event class	Regimentable?
forbid <code>net_egress</code>	controllable	yes (regiment)
forbid <code>git_push</code> (incl. force-push)	controllable	yes (regiment)
forbid <code>fs_write</code> (out-of-scope write)	controllable	yes (regiment)
forbid <code>exec_tool</code>	controllable	yes (regiment)
forbid <code>spawn_child</code>	controllable	yes (regiment)
forbid <code>model_emit_token</code>	uncontrollable	no (detect only)
forbid <code>in_context_read</code>	uncontrollable	no (detect only)
forbid <code>internal_plan</code> (“confident lie”)	uncontrollable	no (detect only)
compound: no <code>net_egress</code> after secret <code>in_context_read</code>	mixed	yes (regiment)

All nine verdicts are the checker’s output, verbatim in content [internal, `b3_controllability.py`]; the eight single-event rows also follow by hand from the theorem, which is the right way to first believe the table. *Numbers by hand*: take “forbid `git_push`.” The policy automaton has one state with every event except `git_push` self-looping; `git_push` $\in \Sigma_c$, so the only event the policy ever disables is controllable, the premise of a controllability violation ($\sigma \in \Sigma_u$ enabled in plant, disabled in policy) is unsatisfiable, and K is controllable — regimentable. Now take “forbid `internal_plan`”: the same one-state automaton disables `internal_plan` $\in \Sigma_u$, the universal plant enables it at the start state, so the empty history $s = \varepsilon \in \bar{K}$ and $\sigma = \text{internal_plan}$ witness $\bar{K}\Sigma_u \cap \bar{L} \not\subseteq \bar{K}$ — detect-only, forever, for any runtime with this alphabet split. *Now you try*: classify “forbid `api_call` after `spawn_child`” and “forbid `internal_plan` after `fs_write`.” (Answers: the first gates a controllable event on a controllable trigger — regimentable; the second must disable an uncontrollable event in its post-trigger state, and the plant enables it there — detect-only.)

This table is the manuscript’s needed re-grading of governance rules, stated correctly: *force-push, unauthorized egress, and out-of-scope writes are matters of structural authority and shall be prevented; a confident lie in a status report is semantic, uncontrollable, and shall be detected*

¹skills/harbor-results/scripts/b3_controllability.py; deterministic, no seed needed; exits 0.

nine-policy classification	
policy	verdict
forbid fs_write	regiment
forbid net_egress	regiment
forbid exec_tool	regiment
forbid git_push	regiment
forbid spawn_child	regiment
compound: no	regiment ★
net_egress after	
in_context_read	
the divider carries the criterion: $\overline{K}\Sigma_u \cap \overline{L} \subseteq \overline{K}$ (Ramadge–Wonham 1987) — no uncontrollable event exits the spec	
forbid	detect-only
model_emit_token	
forbid	detect-only
in_context_read	
forbid	detect-only
internal_plan	

compound case: how it becomes regimentable




Figure 2: The regime diagram. Above the divider, the regimentable regime (Ramadge–Wonham controllable): forbid fs_write / net_egress / exec_tool / git_push / spawn, and the compound “no egress after secret read.” Below, detect-only forever: forbid token emission / in-context read / internal plan / confident falsehood. The divider carries the criterion $\overline{K}\Sigma_u \cap \overline{L} \subseteq \overline{K}$ — no uncontrollable event exits the specification. Classifications reproduce from the product-automaton checker [internal, b3_controllability.py].

and compensated — by audit, reputation, and slashing, which is the adjacent paper’s machinery, not this one’s. Figure 2 draws the same line as a single divider: everything above it gates at the mediation boundary, everything below it is detect-only forever, and the criterion on the divider is the whole theorem in one inequality.

6 The compound case: gate the channel, never the token

The last table row is the one the whole clean-room product line rests on, and it is the case where untrained intuition most often guesses wrong. The policy “no network egress after an in-context read of a secret” mentions an uncontrollable event (in_context_read) — so a naive reading of the single-event rows predicts detect-only. The checker says otherwise:

COMPOUND POLICY: 'no net_egress after in_context_read of a secret'
regimentable? YES

[internal, b3_controllability.py]. The mechanism is visible in the policy automaton: two states, ok and tainted. The uncontrollable read is enabled in both — the policy never tries to forbid it; it lets the read happen and records the taint transition. What the tainted state disables is net_egress alone, and egress is controllable. Check the criterion: every uncontrollable event enabled by the plant is enabled by the policy in every reachable state, so $\overline{K}\Sigma_u \cap \overline{L} \subseteq \overline{K}$ holds and the theorem’s constructive direction hands us the supervisor — track taint, gate the channel.

Three consequences, in increasing order of consequence.

First, the design rule. A compound trigger→effect policy is regimentable iff its effect events are controllable; the trigger may be as uncontrollable as it likes, provided the policy observes

rather than forbids it. Uncontrollable events in the policy’s *condition* are free; uncontrollable events in its *prohibition* are fatal.

Second, the clean room. The Sealed Harbor design (Paper 4) puts an agent alone in a room with a customer’s secret data and promises the secret does not leave. The theorem dictates the only sound architecture: it is impossible to prevent the model from *reading* the secret (the read is uncontrollable — and useless to forbid, since reading is the job), and it is impossible to prevent the model from *incorporating* the secret into tokens and plans (uncontrollable again). What is possible — exactly and only — is to gate every controllable egress channel out of the room, with taint recorded from the moment of exposure. Gate the channel, never the token. Paper 4’s companion noninterference result proves from the other direction that the gate suffices: observations outside the room are identical across secrets except at declared gate releases. Controllability says the gate is the only place a guarantee *can* live; noninterference says a guarantee *does* live there. This controllability boundary is thus the load-bearing assumption behind the clean-room ledger’s claims elsewhere in the program: “all spend is recorded” is precisely complete mediation of Σ_c .

Third, the special case that explains a folk theorem. Product reviews of agent governance converged on the slogan “a rule is enforceable iff it is a pure deterministic function of committed local DB state.” That slogan is our theorem’s special case $\Sigma_c = \{\text{DB writes}\}$: when the only mediated channel is the database commit, the regimentable policies are exactly those expressible over committed state. The general statement strictly dominates it — the slogan cannot express why force-push is preventable by a runtime that also owns the git credential channel, nor why a confident lie stays unpreventable no matter how much state is committed. Widening Σ_c (owning more channels) is the *only* way to grow the regimentable set; no policy-language cleverness does.

7 A synthesis case study: the sandboxed review job

The theorem’s constructive direction is not just an existence proof — it is the synthesis algorithm the Supremica / TCT tool family implements, and it is small enough here to run by hand. Consider a realistic composite policy for a sandboxed code-review job:

- P_1 : no `net_egress`, ever (sandboxed);
- P_2 : no `git_push`, ever (review jobs propose, humans land);
- P_3 : after any `in_context_read` of customer material, no `fs_write` (findings leave through the mediated report channel, not the workspace).

Each P_i is a safety automaton; the composite policy is their synchronous product. P_1 and P_2 are one-state automata, P_3 is the two-state taint automaton of §6, so the product spec has $1 \times 1 \times 2 = 2$ states, `ok` and `tainted`, and the product with the universal one-state plant has the same two reachable states. The controllability check visits both states and tests each of the four uncontrollable events in each: all eight (state, event) pairs are enabled by the spec — neither state ever disables a member of Σ_u — so there are 0 violations and the composite is controllable [verified: eight cases, checkable by eye]. The maximally permissive supervisor falls out of the proof’s construction as a two-row disablement map:

Supervisor state	Disabled ($\subseteq \Sigma_c$)	Everything else
<code>ok</code>	<code>{net_egress, git_push}</code>	enabled
<code>tainted</code>	<code>{net_egress, git_push, fs_write}</code>	enabled

Two states, one observed transition (`in_context_read` flips `ok` to `tainted`), three gates. That table *is* the synthesized supervisor — it is also, not coincidentally, exactly the policy engine a competent engineer would have written by hand. The value of synthesis is not that the answer is surprising; it is that the answer arrives with a maximal-permissiveness certificate (no legal

behavior was sacrificed) and a controllability certificate (no uncontrollable event can force a violation), neither of which hand-written policy engines carry.

Synthesis is equally informative when it fails. Ask it to weaken the detect-only policy “forbid `internal_plan`” to its best preventable approximation: the supremal controllable sublanguage. Let $K' \subseteq K$ be any nonempty controllable sublanguage; it contains ε , the plant enables `internal_plan` at ε , so controllability forces $\varepsilon \cdot \text{internal_plan} \in \overline{K'} \subseteq \overline{K}$ — contradicting that K forbids it. Hence $\sup \mathcal{C}(K) = \emptyset$ [verified: three lines, above]. The theory refuses to pretend: there is no nontrivial preventable weakening of a thought ban — the nearest preventable policy prevents nothing — and any governance document claiming to “partially enforce” such a rule by prevention is describing detection with optimistic vocabulary.

8 Partial observation: the next boundary (future work)

Our theorem assumes the supervisor *observes* every event, even the ones it cannot disable — the daemon logs tokens after emission, reads after completion. Realistic deployments break this in a second place: some events are not only unrefusable but *unobservable* (activations never surfaced; state internal to a third-party tool; another principal’s context). Supervisory control under partial observation is also solved territory: Lin–Wonham [3] characterize enforceability as controllability *plus observability* — the supervisor’s decisions must be constant on observation-equivalence classes of histories, i.e. the policy may not demand different gate decisions after histories the runtime cannot tell apart. Observability strictly shrinks the regimentable set: a policy like “no egress after the model *considered* exfiltration” references unmediated internal *state* with no observable taint transition, so even though its prohibition (egress) is controllable, no supervisor can implement the gate condition. The practical reading: a compound policy is deployable only if its trigger is *witnessed at the mediation boundary* (the secret arrives via a mediated `in_context_read` the daemon performs, as in §6) — taint must be an event the daemon sees, not a state it infers. Mapping the agent-governance policy corpus across the observability line, and connecting it to the assurance-level ladder (which integration tiers witness which triggers), is the natural sequel; we state it as future work rather than claim it, because the interesting part — which real triggers are witnessed at which assurance levels — is an empirical systems question, not a corollary.

9 Related work

Supervisory control of discrete-event systems originates with Ramadge and Wonham [1, 2]; we import controllability, the supremal controllable sublanguage, and the maximally permissive supervisor unchanged, and claim no new control theory. Lin and Wonham [3] supply the partial-observation refinement sketched in §8. Schneider [4] characterizes EM-enforceable policies as safety properties; our contribution relative to Schneider is the refinement under incomplete mediation ($\Sigma_u \neq \emptyset$), where safety is necessary but controllability is the exact boundary. The reference-monitor lineage — Anderson’s complete-mediation obligation [5], Rushby’s kernel-verification standard [6], and Goguen–Meseguer’s policy-as-noninterference framing [7] — supplies the architecture our Σ_c presumes and the companion theory Paper 4 builds on. Recent agent-guardrail work (2024–2026) invokes this security canon for LLM runtimes but, to our knowledge from an August-2026 survey, does not state the controllability characterization; the mapping of agent-runtime enforcement onto supervisory control, and the resulting exact prevented/detected boundary, is the contribution — with the honest corollary that decades of synthesis results (supremal sublanguages, modular and hierarchical supervision, Supremica-class tooling) now import wholesale into agent governance. A final lit-sweep before submission is owed: “not found” is not “proven nonexistent.”

10 Honest boundary

What this theorem does NOT say. (i) It does not say detect-only policies are unenforceable or hopeless — detection feeding audit, reputation, and bond-slashing (Paper 3) is real enforcement; the theorem only says prevention is off the table, so pricing and compensation must be sized for violations that *will* occur. (ii) It does not certify any implementation: the theorem’s Σ_c is the set of events *actually* mediated, and every unenumerated side channel (raw sockets, DNS, inherited file descriptors, /proc, git credential helpers, package managers, clipboard, provider callbacks) silently moves events from Σ_c to Σ_u and shrinks the regimentable set behind the operator’s back — complete mediation of the claimed alphabet is an architectural obligation (Anderson; Rushby) that a red-team, not this proof, discharges. (iii) The classification is relative to the alphabet split: a runtime that gates model steps (e.g. token-level filtering with the model inside the boundary) changes Σ_u and re-grades the table — the theorem is a functor from architectures to boundaries, not one fixed verdict. (iv) The full-observation assumption is load-bearing: policies whose triggers are unwitnessed internal state fall to the Lin–Wonham refinement (§8) even when their prohibitions are controllable. (v) The checker verifies the stated policy automata over the stated alphabet [internal, `b3_controllability.py`]; it is a faithful implementation of the test, not a mechanized proof of the theorem — an Isabelle/Coq formalization remains the submission gate. (vi) Prevention governs *effects*, never *semantics*: no widening of Σ_c ever makes a confident lie preventable, because the lie is constituted, not merely caused, by uncontrollable events.

Reproducibility. The policy table and the compound-case verdict regenerate deterministically from `skills/harbor-results/scripts/b3_controllability.py` (pure Python, no dependencies, no randomness); the two figures regenerate from `docs/harbor-research/figures/src/r5_figures.py`. Every number and verdict in this paper is tagged [verified] (externally recomputable from the cited literature or by hand) or [internal] (regenerates from the named script).

References

- [1] P. J. Ramadge and W. M. Wonham. Supervisory control of a class of discrete event processes. *SIAM Journal on Control and Optimization*, 25(1):206–230, 1987.
- [2] P. J. Ramadge and W. M. Wonham. The control of discrete event systems. *Proceedings of the IEEE*, 77(1):81–98, 1989.
- [3] F. Lin and W. M. Wonham. On observability of discrete-event systems. *Information Sciences*, 44(3):173–198, 1988.
- [4] F. B. Schneider. Enforceable security policies. *ACM Transactions on Information and System Security*, 3(1):30–50, 2000.
- [5] J. P. Anderson. *Computer Security Technology Planning Study*. Technical Report ESD-TR-73-51, U.S. Air Force Electronic Systems Division, 1972.
- [6] J. Rushby. *Noninterference, Transitivity, and Channel-Control Security Policies*. Technical Report CSL-92-02, SRI International, December 1992.
- [7] J. A. Goguen and J. Meseguer. Security policies and security models. In *Proceedings of the IEEE Symposium on Security and Privacy (Oakland)*, pages 11–20, 1982.