

Chapter 0: Foundations and Prerequisites

Pre-LLM Multi-Agent Theory, Transformer Mechanics, and the Kernel Gap

The Authors

September 2026

Express lane: Before constructing an operating kernel for autonomous systems, we must establish the physical and theoretical ground upon which modern AI agents stand. This chapter establishes eight core foundations assumed by the subsequent eight chapters: (i) the dual lineage of autonomous agency connecting classical symbolic multi-agent systems to generative connectionism; (ii) classical multi-agent theory, spanning Bratman practical rationality, Rao–Georgeff BDI_{CTL}^ temporal logic, AgentSpeak(L) operational semantics, normative deontic reasoning, Hewitt–Agha actor concurrency versus Hoare CSP, FIPA-00037 speech acts, Smith’s Contract Net Protocol, Big Brother dynamic epistemic logic, and game-theoretic mechanism design; (iii) modern coding agent architectures, contrasting leading systems and formalizing closed-loop POMDP deliberation (ReAct, Tree of Thoughts, Reflexion); (iv) cognitive control, tool-call SFT, PRMs, and test-time compute; (v) transformer mechanics, KV-cache memory scaling, and context economics; (vi) the tool execution hypervisor, MCP protocols, skill progressive disclosure, and Ramadge–Wonham supervisory control hooks; (vii) multi-agent swarm topologies, durable identity versus ephemeral leases, and zombie salvage; and (viii) the kernel gap, proving mathematically why generative models cannot act as their own operating system.*

1 Introduction: The Unit of Account and the Dual Lineage

Agentic software is usually measured by the agent: capability, runtime, tool access, and parallel copies. That is the wrong unit of account. A process may disappear, rename itself, summarize its mistakes, or hand work off. The durable object must be the work: its authorization, permitted effects, evidence, open obligations, and liability for the outcome.

1.1 The Dual Lineage of Autonomous Agency

Autonomous software agents did not begin with large language models. The problem of coordinating multiple independent computational processes, each pursuing distinct local goals under incomplete information and bounded resources, has occupied distributed systems and artificial intelligence research for over four decades.

When a modern coding agent inspects a git repository, evaluates compiler diagnostics, plans a refactoring sequence, and invokes a bash shell, it operates at the intersection of two previously separate research lineages:

1. **Symbolic Multi-Agent Systems (1975–2010):** Mathematical formalisms for practical rationality, intention persistence, asynchronous message passing, contract negotiation, epistemic logic, and social norms.
2. **Statistical Connectionism (2017–2026):** Deep autoregressive sequence prediction trained over planet-scale code and natural language corpora, capable of zero-shot semantic reasoning, code synthesis, and tool parameterization.

The tension between these lineages defines the subject of this textbook. Modern frontier models supply unprecedented semantic reasoning, but possess neither persistent state, monotonic capability boundaries, nor physical execution guarantees. When an ungrounded model encounters an unexpected tool error, its probabilistic nature induces intention drift, hallucinatory claims of success, and concurrency hazards. To scale from an isolated interactive coding assistant to a collaborative, multi-agent autonomous engineering fleet, we must place statistical inference inside an authoritative, mathematically verified structural envelope where the unit of account is the work rather than the ephemeral process.

1.2 Defining the Agent and the Kernel: Mathematical Form and Architectural Necessity

To reason rigorously about coordination failures, we must first establish formal mathematical definitions for the two primary computational objects in this textbook: the **Agent** and the **Operating Kernel**.

1.2.1 What is an Agent?

In contemporary discourse, the term “agent” is frequently conflated with a stateless language model invocation, a prompt template, or a continuous chatbot loop. In this volume, we reserve the term strictly for an embodied cybernetic process exhibiting closed-loop, temporal goal pursuit against an external state space.

Definition 1.1 (Autonomous Agent). *An **Agent** $\mathcal{A}$ is a 6-tuple:*

$$\mathcal{A} = \langle \mathcal{S}, \mathcal{O}, \mathcal{A}_{ct}, \mathcal{T}, \mathcal{M}, \pi \rangle \quad (1)$$

where:

- $\mathcal{S}$ is the external environment state space (the host filesystem hierarchy, git object graph, POSIX process table, open TCP ports, and network routes).
- $\mathcal{O}$ is the discrete observation space, produced by a sensory projection mapping $\omega : \mathcal{S} \rightarrow \mathcal{O}$ (stdout/stderr streams, exit codes, AST representations, directory manifests, and HTTP responses).
- $\mathcal{A}_{ct}$ is the discrete tool-action alphabet (`exec_command`, `read_file`, `write_file`, `acquire_lock`, `emit_signal`).
- $\mathcal{T} : \mathcal{S} \times \mathcal{A}_{ct} \rightarrow \Delta(\mathcal{S})$ represents the non-deterministic transition dynamics of the external environment.
- $\mathcal{M}$ is the internal cognitive and memory state of the agent, comprising its working context buffer, autoregressive key-value (KV) activations, episodic scratchpad, and externalized goal stack.
- $\pi : \mathcal{M} \times \mathcal{O} \rightarrow \Delta(\mathcal{A}_{ct} \times \mathcal{M})$ is the deliberation policy, which maps the current internal memory and incoming observation to a probability distribution over executable actions and successor memory states.

This definition establishes three fundamental separation boundaries:

1. **Agency vs. Next-Token Prediction:** A bare transformer computes only $P_\theta(x_{t+1} \mid x_{\leq t})$. It becomes an agent $\mathcal{A}$ only when coupled to an observation sensor ω , an action execution interface $\mathcal{A}_{ct}$, and a state transition history $\mathcal{T}$.

2. **Internal Memory vs. External Truth:** The cognitive state $\mathcal{M}$ is subjective, fallible, and volatile. It contains what the model *believes* or *recalls*, which may drift arbitrarily from the true physical state $s \in \mathcal{S}$.
3. **Action Execution vs. Intention:** Generating a string representation of a tool call in $\mathcal{M}$ has zero causal force in $\mathcal{S}$ until it passes through an external execution boundary.

1.2.2 What is an Agent Operating Kernel?

Host operating systems (such as Linux, macOS/XNU, or Windows) are designed to mediate physical hardware: CPU instruction cycles, virtual memory page tables, hardware interrupts, and block storage devices. They possess zero semantic awareness of AI agents:

- The host kernel cannot detect when two LLMs are making conflicting semantic edits to overlapping AST nodes across different git worktrees.
- The host kernel cannot determine whether a generated bash script is pursuing a legitimate refactoring or hallucinating nonexistent compiler flags in an infinite retry loop.
- The host kernel cannot enforce capability attenuation across natural language delegation chains, nor can it meter token spend against a financial budget.

Definition 1.2 (Agent Operating Kernel). *An **Agent Operating Kernel** is the minimal, non-bypassable Trusted Computing Base (TCB) operating at the application and multi-agent coordination plane. It provides five non-negotiable architectural guarantees:*

1. **Single-Writer Effect Serialization:** *Mediates all state-mutating side effects across shared resources (git commits, lock tables, port assignments, and workspace directories) via a single-writer write-ahead log (WAL), enforcing happened-before causal order.*
2. **Cryptographic Capability Attenuation:** *Issues unforgeable, key-bound credentials (such as Harbor Cards and macaroons) whose authority and temporal validity can only narrow at each delegation hop.*
3. **Deterministic Resource Accounting:** *Enforces strict, non-negotiable caps on token spend, dollar budgets, execution wall-clock time, and subprocess concurrency, failing closed when thresholds are reached.*
4. **Tamper-Evident Lineage Logging:** *Records every admitted action, sensor observation digest, and authorization receipt in an append-only, content-addressed Merkle event log.*
5. **Supervisory Liveness and Salvage:** *Monitors agent heartbeat telemetry, detects cognitive loops or silent hangs, and reclaims abandoned file claims and port leases via deterministic zombie salvage.*

1.3 Is a Daemon Ever a Kernel? Microkernels, Application Planes, and the TCB

This brings us to the titular question that opens this inquiry: *Can a user-space background daemon legitimately be considered an operating kernel?*

In traditional monolithic operating system design (exemplified by UNIX, Linux, and FreeBSD), the term “kernel” is synonymous with execution in Ring 0 supervisor mode, with direct access to privileged CPU instructions and hardware page tables. By this narrow hardware-privilege definition, a daemon executing in Ring 3 user-space cannot be a kernel.

However, modern operating systems theory—beginning with Per Brinch Hansen’s RC 4000 Nucleus, Jochen Liedtke’s L4 microkernel family, Mach, QNX, and formal capability systems like seL4 and KeyKOS—fundamentally dismantled this conflation between *hardware privilege* and

kernel functionality. In microkernel philosophy, the kernel is defined not by how much hardware it directly manipulates, but by its formal role as the **Trusted Computing Base (TCB)**: the minimal, authoritative mechanism that cannot be bypassed and upon whose correctness the security and integrity of the entire system depend. In an L4 or QNX microkernel, device drivers, virtual memory pagers, and file systems execute entirely in user space as discrete, message-passing server *daemons*.

When reasoning about autonomous AI engineering fleets, the primary risk surface does not lie in x86-64 page faults or PCI bus contention; it lies in *uncoordinated agent side effects, race conditions on shared repositories, unauthenticated delegation, and runaway resource consumption*. An Agent Operating Kernel (such as **agentsd**) sits as a supervisory mediator between untrusted LLM processes and external environments:

1. **Unbypassable Mediation:** When agents execute inside confined operating-system sandboxes (macOS Seatbelt, Linux Landlock/Bubblewrap, or isolated microVMs), their direct access to host network sockets, git repositories, and external APIs is severed. The only egress channel available to the agent is an authenticated UNIX domain socket or local RPC pipe terminating at **agentsd**. Because the agent cannot circumvent this channel to mutate shared reality, **agentsd** satisfies Lampson’s classical reference monitor property: it is *complete, tamper-proof, and verifiable*.
2. **Single-Writer Authority:** The daemon holds exclusive write locks over the persistent state database (SQLite WAL) and the shared worktree index. No agent can unilaterally declare a lock acquired or a commit finalized without the daemon’s serialized commit receipt.
3. **Failure-Closed Invariant Enforcement:** If the supervisory daemon halts, crashes, or rejects an invalid capability card, all external mutation ceases immediately. The fleet fails closed rather than failing open.
4. **Subprocess Independence:** An agent running inside a container or Python runtime may suffer out-of-memory errors, syntax panics, or hallucinated token loops. The supervisory daemon runs as an independent OS service (supervised by **launchd** or **systemd**), detached from the lifecycles of the agents it governs.

Therefore, functionally, mathematically, and architecturally, a persistent supervisory daemon that acts as the single-writer mediator of state, capability attenuation, and resource accounting **is an operating kernel at the application and coordination plane**.

1.4 Eight Propositions of Accountable Coordination

This book develops that claim in eight propositions, one per chapter, in the order the book is read:

1. **Prevention begins at the effect boundary (The Single-Writer Kernel).** Rules can be enforced before an act only when the system controls the channel through which the act occurs. Local effects therefore pass through one serial authority; semantic failures that no gate can decide remain matters of detection and remedy.
2. **Delegation only narrows (The Anchor Protocol).** Every child capability is a strict subset of its parent’s authority and lifetime. Authentication says who signed; attenuation says what the signature may authorize. Both are needed at every hop.
3. **Confinement is a property of the room, not of the tenant (The Sealed Harbor).** Two owners who trust neither each other nor the venue can still take one answer out of a sealed room when its only exit is a metered slot: silence except through the slot, a gate on the only enforceable boundary, a budget that conserves, and canaries for what the gate

cannot prove. What the room cannot promise is stated as a limit, at the same prominence as the claims.

4. **Supervision has an information cost (Attention and Supervision).** A digest cannot guarantee that a human will notice any one of many relevant events unless it spends enough bits, attention, or artifact openings to distinguish them. A summary is therefore an index into evidence, not a substitute for it.
5. **Accountability must outlive the process (Spawn-to-Person).** A task acquires a durable principal, capability history, evidence record, obligations, and outcome. Restarting an agent or choosing a new name must not erase that case file or refill the resources it already consumed.
6. **Exchange requires separated instruments (Agent Transactions).** Payment rewards the requested work. Collateral prices a bounded failure. Evidence supports a claim. Settlement closes it exactly once. Combining those functions in one score or token makes each of them less trustworthy.
7. **Uncertainty needs an institution (The Harbor Economy).** Some harms cannot be prevented and some outcomes cannot be judged mechanically. Witnesses, randomized audits, bounded bonds, appeals, and conservation rules turn that residual uncertainty into a process that can be inspected and contested.
8. **Federation relays evidence, not sovereignty (Federated Harbors).** Mutually distrustful machines may exchange signed capabilities, revocations, witness roots, and settlement claims without sharing a database or accepting a global ruler. Transport can carry evidence; it cannot decide what either machine is permitted to do with it.

Read these as a dependency chain: each chapter uses earlier mechanisms and states what they cannot guarantee.

1.5 What Counts as Knowing: Epistemic Bounds and Evidence Channels

Every consequential statement in this volume has a status and a boundary. Derivations establish relations under named assumptions. Finite checks exhaust a stated state space, not every deployment. Simulations estimate behavior under a stated generator; repository tests exercise named paths; running code proves availability, not universal mediation. Counterexamples remain in the record because they mark claims the research killed rather than conclusions it merely failed to prove.

Evidence is typed as well. When a chapter reports what a harbor observed, it names the channel that produced the observation, and each channel has a blind spot that no amount of volume removes:

- A *deterministic verifier* (a test, a type checker, a model checker) shows that a stated property holds of the artifact it examined, and nothing about properties it did not state.
- A *provenance witness* (a signature, a hash, a Merkle inclusion, an append-only log) shows attribution and tamper-evidence relative to a commitment, never that every event was logged or that a description is true.
- A *protected meter* (a counter its subject cannot rewrite) shows a quantity, not its cause.
- *Expert judgment* shows an assessment under a named rubric and carries the assessor's error rate.
- *Buyer preference* shows what was chosen, not that the choice was right.
- A *delayed outcome* shows what happened, and arrives too late to prevent it.

The reader should ask the same question of any dashboard: which channel, and what can it not show.

?? gives the implementation boundary; ?? records the results and their limits. The companion `docs/harbor-research/mega-volume-epistemic-manifest.yaml` preserves each result’s source, status, conditions, boundary, chapter home, and verification artifact. The chapters stand alone; the book’s claim is the chain they make.

2 Classical Multi-Agent Systems Foundations

Before analyzing neural coding agents, we formalize the foundational frameworks developed during the classical era of multi-agent systems.

2.1 Intentional Agency and Practical Rationality

Classical decision theory models an agent as an expected utility maximizer balancing subjective beliefs $\mathcal{B}$ and desires $\mathcal{D}$. In his seminal 1987 work *Intention, Plans, and Practical Reason*, philosopher Michael Bratman demonstrated that this two-component model is computationally inadequate for resource-bounded agents operating in dynamic environments [3].

A resource-bounded agent cannot continuously compute optimal policies from first principles at every millisecond. Practical rationality requires a third mental attitude: **Intention** ($\mathcal{I}$). Intentions are not merely desires with high utility; they represent *conduct-controlling pro-attitudes to which the agent has actively committed*.

Bratman established four defining functional roles of intentions:

1. **Drive Means-End Reasoning:** An intention to achieve goal ϕ obligates the agent to formulate sub-plans and allocate computational cycles until ϕ is satisfied.
2. **Filter of Admissibility:** Once an intention ϕ is adopted, candidate options inconsistent with ϕ (given current beliefs $\mathcal{B}$) are ruled inadmissible without evaluating their cost-benefit payoff.
3. **Stability and Commitment:** Intentions resist reconsideration in the absence of relevant environmental interrupts or belief updates, preventing computational paralysis.
4. **Coordination Scaffolding:** Intentions enable interpersonal coordination: other agents can safely form expectations regarding an agent’s future actions based on its declared commitments.

2.2 The Rao–Georgeff BDI Temporal Logic (BDI_{CTL^*})

Rao and Georgeff [7, 8] formalized Bratman’s theory into a multimodal branching-time temporal logic (BDI_{CTL^*}). A temporal world model $\mathcal{M}$ is defined as an 8-tuple:

$$\mathcal{M} = \langle W, T, \prec, U, \mathcal{B}, \mathcal{D}, \mathcal{I}, \Phi \rangle \quad (2)$$

where W is a set of possible worlds, T is a discrete set of time points ordered by tree-branching relation $\prec$, U is the universe of discourse, Φ interprets predicates across worlds and times, and $\mathcal{B}, \mathcal{D}, \mathcal{I} \subseteq W \times T \times W$ are accessibility relations for Beliefs, Desires, and Intentions.

For world $w \in W$ at time $t \in T$:

$$\mathcal{B}_t(w) = \{w' \in W \mid (w, t, w') \in \mathcal{B}\} \quad (3)$$

$$\mathcal{D}_t(w) = \{w' \in W \mid (w, t, w') \in \mathcal{D}\} \quad (4)$$

$$\mathcal{I}_t(w) = \{w' \in W \mid (w, t, w') \in \mathcal{I}\} \quad (5)$$

Semantics of the mental modalities at state $\langle w, t \rangle$:

$$\mathcal{M}, w, t \models \text{Bel}(\phi) \iff \forall w' \in \mathcal{B}_t(w), \mathcal{M}, w', t \models \phi \quad (6)$$

$$\mathcal{M}, w, t \models \text{Des}(\phi) \iff \forall w' \in \mathcal{D}_t(w), \mathcal{M}, w', t \models \phi \quad (7)$$

$$\mathcal{M}, w, t \models \text{Int}(\phi) \iff \forall w' \in \mathcal{I}_t(w), \mathcal{M}, w', t \models \phi \quad (8)$$

Beliefs satisfy the *KD45* axiom system (deductive closure **K**, consistency **D**, positive introspection **4**, and negative introspection **5**). Desires and Intentions satisfy seriality (**D**), ensuring an agent never desires or intends logical contradictions ($\neg\text{Des}(\perp)$ and $\neg\text{Int}(\perp)$).

2.2.1 Compatibility and Weak Realism

In rational agency, an agent's desires and intentions must be constrained by its beliefs. Under the **Weak Realism** axiom, an agent cannot intend what it believes to be impossible:

$$\forall w \in W, \mathcal{I}_t(w) \cap \mathcal{B}_t(w) \neq \emptyset \quad (9)$$

Syntactically, for temporal achievement goals:

$$\text{Int}(\mathbf{A}\Diamond\phi) \rightarrow \neg\text{Bel}(\neg\mathbf{E}\Diamond\phi) \quad (10)$$

where $\mathbf{A}\Diamond\phi$ denotes “ ϕ holds eventually on all future paths,” and $\mathbf{E}\Diamond\phi$ denotes “ ϕ holds on at least one branch.”

2.3 Formal Commitment Strategies

A commitment strategy dictates the precise formal conditions under which an agent drops an intention ϕ . Rao and Georgeff [7] formalized three canonical strategies:

Commitment Strategy	Formal Drop Axiom
Blind (Fanatical)	Maintains $\text{Int}(\mathbf{A}\Diamond\phi)$ until $\text{Bel}(\phi)$ holds. $\text{Int}(\mathbf{A}\Diamond\phi) \wedge \neg\text{Bel}(\phi) \rightarrow \mathbf{A}(\text{Int}(\mathbf{A}\Diamond\phi) \mathcal{U} \text{Bel}(\phi))$. <i>Pathology:</i> Induces infinite loops if ϕ becomes physically impossible.
Single-Minded	Drops ϕ if achieved <i>or</i> believed impossible. $\text{Int}(\mathbf{A}\Diamond\phi) \wedge \neg\text{Bel}(\phi) \wedge \text{Bel}(\mathbf{E}\Diamond\phi) \rightarrow \mathbf{A}(\text{Int}(\mathbf{A}\Diamond\phi) \mathcal{U} (\text{Bel}(\phi) \vee \neg\text{Bel}(\mathbf{E}\Diamond\phi)))$. <i>Standard:</i> Balances persistence with dynamic feasibility.
Open-Minded	Drops ϕ if achieved, impossible, <i>or no longer desired</i> . $\text{Int}(\mathbf{A}\Diamond\phi) \wedge \dots \wedge \text{Des}(\mathbf{A}\Diamond\phi) \rightarrow \mathbf{A}(\text{Int}(\mathbf{A}\Diamond\phi) \mathcal{U} (\text{Bel}(\phi) \vee \neg\text{Bel}(\mathbf{E}\Diamond\phi) \vee \neg\text{Des}(\mathbf{A}\Diamond\phi)))$.

Table 1: Canonical commitment strategies in BDI agent architectures.

Design Invariant (Commitment Persistence).

An agent's intention $\iota \in \mathcal{I}$ must not be abandoned merely because a single step incurs a transient tool failure. Intention reconsideration is an explicit state transition triggered by invariant failure or objective discharge, not an accidental consequence of context window truncation.

2.4 Operational Semantics of AgentSpeak(L) and Normative Extensions

AgentSpeak(L) [?] provides structural operational semantics (SOS) for executing BDI agents. An agent configuration is a tuple $C = \langle bs, ps, is, \mathcal{E}, A \rangle$, where bs is the belief base, ps is the plan library, is is the set of active intention stacks, $\mathcal{E}$ is the event queue, and A is the action history.

The operational execution cycle is governed by three deterministic or priority-driven selection functions:

1. $S_E(\mathcal{E}) = \langle te, i \rangle$: Selects an event from $\mathcal{E}$ (e.g., belief change $+b$, goal addition $+!g$).
2. $S_O(\text{App})$: Selects a plan (p, θ) from the applicable plans whose context is entailed by $bs\theta$.
3. $S_I(is) = i$: Selects which intention stack advances its execution step in the current cycle.

When an action a in plan p fails during host execution, AgentSpeak semantics do not silently abort the agent; instead, the failure transition posts a failure event $-!g$ to $\mathcal{E}$, popping the failed plan frame and invoking explicit recovery plans:

$$\frac{\epsilon = \langle te, i \rangle, \quad \text{App}(\text{Rel}(te, ps), bs) = \emptyset}{\langle bs, ps, is, \mathcal{E}, A \rangle \longrightarrow \langle bs, ps, is, \mathcal{E} \cup \{\langle -!g, \text{pop}(i) \rangle\}, A \rangle} \quad (11)$$

Figure ?? synthesizes this operational choice architecture, illustrating how autonomous systems select between reactive rules, static classical planning, and BDI execution, and how S_E , S_O , and S_I arbitrate runtime goals.




Figure 1: BDI execution flow and operational choice architecture. Adapted from AgentSpeak(L) semantics [?], distinguishing reactive rules, static STRIPS-style planners, and dynamic BDI agents. The three core selection functions (S_E for events, S_O for applicable plans, and S_I for intention stacks) govern runtime scheduling, with explicit plan failure handling via the $-!g$ recovery event.

2.4.1 Normative Extensions: Deontic Constraints and Regret Minimization

In institutional multi-agent systems, agents are constrained by external **norms** (deontic logic). A norm is defined as a 6-tuple $N = \langle \text{id}, \text{Type}, \phi_{\text{act}}, \phi_{\text{tgt}}, d, \text{Sanction} \rangle$, where $\text{Type} \in \{\mathcal{O}, \mathcal{F}, \mathcal{P}\}$ (Obligation, Prohibition, Permission), ϕ_{act} is the activation context, ϕ_{tgt} is the target state, d is the deadline, and Sanction is the penalty for violation.

To reconcile institutional norms with internal intentions, agents maintain an **Abstract Norm Base** (ANB) and evaluate instantiated norms in the **Norm Instance Base** (NIB).

When norm instances conflict with active goals ($Plans(is \cup \mathcal{N}) = \emptyset$), the agent resolves the normative conflict by partitioning candidate norms into **Maximal Non-Conflicting Subsets** (MNCS), selecting the subset that minimizes maximum penalty under a **minimax regret** criterion:

$$S^* = \arg \min_{S \in \Omega} \max_{n_v \in \mathcal{N} \setminus S} \text{Cost}(\text{Sanction}(n_v)) \quad (12)$$

This ensures that any deliberate norm violation is recorded in an immutable ledger alongside its formal justification.

2.5 Concurrency: Hewitt–Agha Actor Model vs. Hoare CSP

Multi-agent coordination requires formal concurrency models. Classical computer science produced two foundational, mathematically rigorous paradigms for message-passing concurrency:

Dimension	Hewitt–Agha Actor Model [5, 1]	Hoare’s CSP (1978) [?]
Primitive	Autonomous Actor with private state	Sequential Process and Named Channel
Synchronization	Asynchronous (non-blocking send)	Synchronous (rendezvous handshake)
Addressing	First-class Mailbox addresses	Explicit Channel identifiers
Buffering	Unbounded mailbox queue	Zero buffer (unbuffered channels)
State Evolution	Replacement behavior ($\text{become}(B')$)	Sequential instruction trace
Failure Mode	Mailbox overflow; unhandled messages	Circular rendezvous deadlock
Modern Descendant	Erlang/OTP, Akka, Ray, agentsd	Go channels, Rust crossbeam

Table 2: Message-passing concurrency: Hewitt–Agha actors versus Hoare CSP.

In the Actor model, an actor processes one message from its mailbox by executing a bounded 3-tuple:

$$\text{Actor}(S, m) \xrightarrow{\text{atomic}} (S', \{m_1, \dots, m_k\}, \{\alpha_{\text{new},1}, \dots, \alpha_{\text{new},j}\}) \quad (13)$$

where S' is the designated replacement behavior, $\{m_k\}$ are asynchronous outbound messages, and $\{\alpha_{\text{new}}\}$ are newly spawned actors. This share-nothing encapsulation provides the blueprint for sandboxed agent processes.

Figure ?? diagrams the Hewitt–Agha actor concurrency model, highlighting its private state encapsulation and asynchronous mailbox versus Hoare’s synchronous CSP channel rendezvous.




Figure 2: The Hewitt–Agha Actor Concurrency Model versus Hoare’s Communicating Sequential Processes (CSP). An actor α encapsulates private state S_t and processes messages from its FIFO mailbox by atomically executing exactly three primitives: sending asynchronous messages, spawning new actors, and designating its successor replacement behavior (**become**).

2.6 Communicative Action and Speech Acts (FIPA-00037)

Standard computer networks exchange raw data packets; autonomous agents exchange **speech acts**. Grounded in Austin’s distinction between locutionary (syntactic utterance), illocutionary (communicative intent), and perlocutionary (consequential effect) acts [2], and Searle’s taxonomy of illocutionary forces [9], the Foundation for Intelligent Physical Agents standardized the **FIPA ACL Communicative Act Library** (FIPA-00037) [4].

Every communicative act $a(s, r, \phi)$ between sender s and receiver r is formally defined by its **Feasibility Preconditions** (FP) and intended **Rational Effect** (RE):

Act	Category	Feasibility Precondition (FP)	Rational Effect (RE)
inform	Assertive	$Bel_s(\phi) \wedge \neg Bel_s(Bel_r(\phi))$	$Bel_r(\phi)$
request	Directive	$FP(\alpha)[s \setminus r] \wedge Bel_s(\text{Agent}(r, \alpha))$	$\text{Done}(\alpha)$
query-if	Directive	$\neg Bel_s(\phi) \wedge \neg Bel_s(\neg \phi)$	$Bel_s(Bel_r(\phi) \vee Bel_r(\neg \phi))$
propose	Commissive	$Bel_s(\text{Int}_s(\text{Done}(\alpha, \phi)))$	$Bel_r(\text{Int}_s(\text{Done}(\alpha, \phi)))$
agree	Commissive	$Bel_s(\text{Int}_s(\text{Done}(\alpha)))$	$Bel_r(\text{Int}_s(\text{Done}(\alpha)))$
refuse	Directive	$Bel_s(\neg \text{Int}_s(\text{Done}(\alpha)))$	$Bel_r(\neg \text{Int}_s(\text{Done}(\alpha)))$
failure	Assertive	$Bel_s(\text{Failed}(\alpha, \psi))$	$Bel_r(\text{Failed}(\alpha, \psi))$

Table 3: FIPA-00037 communicative performatives with mental-state preconditions and rational effects.

In an autonomous agent operating kernel, performatives cease to be informal text; they are cryptographically signed, immutable protocol events backed by content-addressed receipts.

Figure ?? formalizes this communicative act transition graph, illustrating how performatives transition across sender proposal, counter-agent acceptance or refusal, and verified execution delivery.




Figure 3: FIPA-00037 Communicative Speech Act state transition machine. Messages are typed performatives governing mental-state preconditions and rational effects across sender and receiver agents.

2.7 Task Allocation: Smith’s Contract Net Protocol (1980)

Reid G. Smith’s (1980) **Contract Net Protocol (CNP)** [?] formalized bilateral market-inspired negotiation between a Manager and distributed Contractor agents:

1. **Task Announcement:** Manager broadcasts an announcement $\langle \text{task_id}, \text{eligibility}, \text{spec}, t_{\text{expire}} \rangle$.
2. **Bidding:** Eligible contractors evaluate local capacity and submit bids $\langle \text{bid_id}, \text{cost}, \text{eta}, \text{quality} \rangle$.
3. **Bid Evaluation and Award:** Manager ranks bids by utility $\mathcal{U}_M(b)$, issuing an **accept-proposal** to the winner and **reject-proposal** to others.

4. **Execution and Receipt:** Contractor executes the task and returns signed execution receipts or formal failure disclosures.

Figure ?? traces the complete execution lifecycle of Smith’s Contract Net Protocol, illustrating the interaction between the coordinating Manager and candidate Contractor nodes.




Figure 4: Smith’s (1980) Contract Net Protocol (CNP) negotiation cycle between a Manager agent and distributed Contractor candidates, establishing competitive task allocation through announcements, bids, contracts, and settlement.

2.8 Dynamic Epistemic Logic and Big Brother Logic

In multi-agent swarms, coordination requires reasoning about higher-order knowledge: what agent i knows about what agent j knows. In modal epistemic logic ($S5_n$), Kripke relations R_i represent epistemic indistinguishability.

Public Announcement Logic (PAL) models communication as epistemic model restriction: when proposition ψ is announced, worlds where ψ is false are eliminated ($W' = \{w \in W \mid \mathcal{M}, w \models \psi\}$), establishing common knowledge $C_G\psi$.

Charrier et al. [?] developed **Big Brother Logic (BBL)**, formalizing epistemic reasoning over embodied agents equipped with directed observation cones (vision sets $V(a)$). If agent a_1 observes agent a_2 ($a_2 \in V(a_1)$), a_1 knows what a_2 observes, yielding $K_{a_1}K_{a_2}(\phi)$. By inverting verification into a planning problem, BBL proves that multi-agent systems can achieve target epistemic goals Φ by autonomously reconfiguring observation cones and network routes.

2.9 Game-Theoretic Foundations and Mechanism Design

When agents have independent objective functions, coordination cannot assume altruism. Shoham and Leyton-Brown [?] systematized the algorithmic and game-theoretic foundations of multi-agent systems.

2.9.1 Social Choice Impossibilities

Arrow’s Impossibility Theorem (1951) and the Gibbard–Satterthwaite Theorem (1973/1975) establish that in ordinal voting environments with three or more alternatives, every strategy-proof (dominant-strategy incentive compatible) social choice rule is necessarily a dictatorship. Consequently, in multi-agent engineering swarms without monetary transfers, strategic manipulation of priority rankings is mathematically unavoidable.

Figure ?? illustrates why democratic preference voting collapses into cyclic deadlocks (the Condorcet paradox) in multi-agent settings, and how the operating kernel avoids social choice pathology by evaluating candidate state transitions through deterministic attestation gates rather than voting.




Figure 5: Arrow’s Impossibility Theorem in multi-agent swarms versus the Kernel Attestation Gate. Direct preference voting over candidate actions induces intransitive Condorcet cycles ($A \succ B \succ C \succ A$). The operating kernel eliminates subjective voting by gating state transitions through deterministic attestation, mutual exclusion locks, and market pricing.

2.9.2 The VCG Mechanism and Dominant-Strategy Incentive Compatibility

In quasi-linear environments where agents have private valuations $v_i(o, \theta_i)$ and make monetary transfers p_i , the **Vickrey–Clarke–Groves (VCG)** mechanism aligns private incentives with global social welfare $W(o) = \sum_i v_i(o, \theta_i)$.

Under VCG, the mechanism determines the socially optimal outcome:

$$o^* = \arg \max_{o \in O} \sum_j v_j(o, \hat{\theta}_j) \quad (14)$$

and charges each agent i its Clarke pivot payment:

$$p_i(\hat{\theta}) = \max_{o \in O} \sum_{j \neq i} v_j(o, \hat{\theta}_j) - \sum_{j \neq i} v_j(o^*, \hat{\theta}_j) \quad (15)$$

The Clarke payment equals the exact negative externality agent i ’s presence imposes on all other participants. Under VCG, truth-telling ($\hat{\theta}_i = \theta_i$) is a strictly dominant strategy (DSIC).

2.9.3 The Myerson–Satterthwaite Boundary

The Myerson–Satterthwaite Theorem (1983) proves that no bilateral trade mechanism can simultaneously achieve: (1) Dominant-strategy incentive compatibility; (2) Ex-post Pareto efficiency; (3) Individual rationality; and (4) Budget balance without an external subsidy. In local agent economies (Chapter 6), resource clearing must explicitly trade off budget balance against exact allocative efficiency.

3 Modern Coding Agent Architectures

We now transition from classical multi-agent foundations to contemporary coding agents, examining how modern LLM-based harnesses interface with host execution environments.

3.1 The Contemporary Landscape: Systems Comparison

Contemporary coding agents bridge unstructured natural language reasoning with deterministic host tools (shells, AST parsers, compilers, git checkouts). Table ?? compares the five leading production architectures:

Agent	Isolation Boundary	Context Strategy	Tool Interface	Rollback Engine
Claude Code	Subprocess + Seatbelt	Progressive disclosure	File edit, bash, MCP	Git commit checkpoint
Codex CLI	Container sandbox	History compaction	Sandboxed patch apply	Ephemeral container reset
Gemini CLI	Local / Cloud Run	1M–2M token window	POSIX CLI, GCP SDK	Workspace rewind
Aider	Host process	Repo-map (Tree-sitter)	Git-backed diff engine	Automatic git commit undo
Devin	MicroVM (Firecracker)	Multi-buffer browser/shell	Full root pseudo-terminal	Snapshot restoration

Table 4: Architectural comparison of contemporary production coding agents.

3.2 Closed-Loop Deliberation vs. Open-Loop Sampling

Stateless code generation evaluates an autoregressive distribution $Y \sim P_\theta(Y \mid X)$ in an open loop. Any syntactic, semantic, or environmental error compounds irreversibly:

$$\text{ErrorRate}_{\text{open-loop}}(T) = 1 - \prod_{t=1}^T (1 - \epsilon_t) \quad (16)$$

In contrast, an agent harness operates as a Partially Observable Markov Decision Process (POMDP):

$$\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \Omega, \mathcal{O}, \gamma \rangle \quad (17)$$

where $\mathcal{S}$ is the true environment state (filesystem, git index, running processes), $\mathcal{A}$ is the action space (tool invocations), Ω is the observation space (stdout, stderr, exit codes), and $\mathcal{O}$ generates observations. The agent maintains a belief state $b_t = \Pr(s_t \mid h_t)$ conditioned on history $h_t = (a_0, o_0, \dots, a_{t-1}, o_{t-1})$.

3.2.1 Deliberation Frameworks

Four primary frameworks govern closed-loop agent reasoning:

1. **ReAct (Reasoning and Acting) [?]**: Interleaves natural language reasoning traces (“Thoughts”) with structured tool executions (“Actions”) and environment feedback (“Observations”). When factual uncertainty is encountered, reasoning halts and an exploratory tool is invoked.
2. **Tree of Thoughts (ToT) [?]**: Explores search trees over candidate thought steps using BFS, DFS, or MCTS, evaluating intermediate node viability via self-critique rubrics and backtracking upon failure.
3. **Reflexion [?]**: Employs verbal reinforcement learning. Upon task failure, the model generates linguistic self-reflections that are stored in an episodic memory buffer and prepended to subsequent execution prompts.
4. **Plan-and-Solve**: Divides execution into a macro-planning phase (generating ordered sub-goals) followed by iterative sub-task execution with explicit replanning interrupts.

3.3 Cognitive Control, Training, and Fine-Tuning

Coding agents are trained across a three-stage continuum:

1. **Pre-Training:** Self-supervised next-token autoregressive modeling over trillions of tokens of source code, version history, documentation, and technical literature:

$$\mathcal{L}_{\text{PT}}(\theta) = -\mathbb{E}_{x \sim \mathcal{D}} \left[\sum_t \log P_\theta(x_t \mid x_{<t}) \right] \quad (18)$$

2. **Tool-Use Supervised Fine-Tuning (SFT):** Training over multi-turn trajectories with *loss masking*: loss gradients are computed strictly over model-generated tool calls and arguments, masking prompt context and tool output tokens.
3. **Reinforcement Learning Alignment (RLHF / DPO):** Direct Preference Optimization optimizes policy weights over paired trajectories without training an explicit reward model:

$$\begin{aligned} \mathcal{L}_{\text{DPO}}(\theta; \theta_{\text{ref}}) = & -\mathbb{E}_{(x, y_w, y_l)} \left[\log \sigma \left(\beta \log \frac{\pi_\theta(y_w \mid x)}{\pi_{\text{ref}}(y_w \mid x)} \right. \right. \\ & \left. \left. - \beta \log \frac{\pi_\theta(y_l \mid x)}{\pi_{\text{ref}}(y_l \mid x)} \right) \right] \end{aligned} \quad (19)$$

3.3.1 Test-Time Compute and Process Reward Models

Reasoning-focused frontier models shift compute to inference time. Instead of sampling greedily, inference search trees explore alternative reasoning chains scored by **Process Reward Models (PRMs)** that evaluate correctness at each logical derivation step:

$$\text{Score}(\tau) = \prod_{t=1}^T \text{PRM}(c_t \mid c_{<t}, x) \quad (20)$$

When an intermediate verification score falls below a threshold τ_{cut} , the search beam backtracks to a previous checkpoint.

3.4 Instruction Hierarchy and Context Economics

A fundamental vulnerability of LLM architectures is the absence of hardware-enforced privilege separation between instructions and data. An agent that reads an untrusted file or issue comment is susceptible to **indirect prompt injection**, where malicious text commandeers model control flow. Production harnesses enforce strict instruction hierarchies: system prompts occupy privileged framing, while user inputs and tool outputs are quarantined in escaped delimitations.

3.4.1 The KV-Cache Asymptotics

During autoregressive decoding, key and value activations for preceding tokens are stored in high-bandwidth memory (HBM). For a model with L layers, N_h KV-heads, dimension d_h , and precision bytes b_p , the memory footprint for a context of length T is:

$$\text{Memory}_{\text{KV}} = 2 \times L \times N_h \times d_h \times T \times b_p \quad \text{bytes} \quad (21)$$

For a 70-billion parameter model ($L = 80, N_h = 8, d_h = 128, b_p = 2$), each context token consumes ≈ 320 KB. At $T = 128,000$ tokens, the KV-cache consumes ≈ 40.96 GB of VRAM per concurrent agent.

Theorem (The Ephemeral Memory Limit).

The context window of a transformer is volatile, expensive, and subject to attention dilution (effective information retrieval degrades with sequence depth [6]). Consequently, long-term system state cannot be maintained by continuously appending tool execution logs to the prompt. Durable state must reside in an external, indexed, and ACID-compliant storage engine.

4 Tools, Protocols, Skills, and Hooks

To interact safely with the host environment, an agent harness must govern tool dispatch, encapsulate domain knowledge, and enforce runtime invariants.

4.1 Tool Dispatch and the Model Context Protocol (MCP)

The Model Context Protocol (MCP) standardizes agent-to-tool communication using JSON-RPC 2.0. MCP defines four core primitives:

- **Tools:** Executable actions with strict JSON-Schema input validation and structured result payloads.
- **Resources:** URI-addressable read-only context (files, database rows, logs).
- **Prompts:** Parametric prompt templates engineered for specific operational workflows.
- **Sampling:** Protocol affordances enabling servers to request nested LLM inference through the client harness under client permission boundaries.

4.2 Skill Architectures and Progressive Disclosure

A critical architectural challenge in agent design is *context bloat*. Injecting complete documentation, code examples, and reference manuals for hundreds of specialized tools into the agent’s system prompt exhausts context limits and induces severe attention dilution [6]. If an agent repository maintains 500 specialized domain skills, each averaging 10,000 tokens of documentation, monolithic injection would require 5,000,000 tokens per inference turn—exceeding the physical memory capacity of frontier models and incurring catastrophic monetary and latency costs.

Modern agent harnesses resolve this dilemma via **Three-Tier Progressive Disclosure** (Figure ??):

1. **Tier 1 (Surface Frontmatter Index):** The system prompt contains only lightweight index entries: the skill’s unique identifier and a concise 1-to-2 sentence description (10^1 – 10^2 tokens per skill). Across a library of 500 skills, this consumes a static footprint of $\approx 15,000$ tokens ($\approx 7.5\%$ of a 200k context window). Because Tier 1 is static across all turns, modern inference providers cache this prompt prefix in HBM, reducing input costs by up to 90%. The probability of inclusion in the prompt is $P = 1.0$.
2. **Tier 2 (On-Demand Instruction Activation):** When the agent’s semantic classifier or planning policy determines that a task requires a specific capability, it invokes a read tool to load the corresponding `SKILL.md` into its working context (10^3 – 10^4 tokens). Tier 2 contains operational instructions, parameter constraints, decision tables, and error recovery protocols. Its activation probability is conditionally sparse ($P \approx 10^{-1}$), loaded only for the duration of the subtask and purged or compacted when the goal is achieved.

3. **Tier 3 (Deep Reference and Subprocess Execution):** Beneath the instruction manual lies the deepest architectural layer: specialized executable scripts (`scripts/`), schemas, and large reference manuals (`references/`), spanning 10^4 – 10^5 tokens. Crucially, the agent *does not load Tier 3 assets into its language context*. Instead, Tier 2 instructs the agent to execute Tier 3 scripts in an isolated host subshell or query specific byte offsets on demand. Because invoking a specialized script or deep reference is needed only for esoteric domain edge-cases, its execution probability is orders of magnitude lower ($P \approx 10^{-3}$ to 10^{-4}).




Figure 6: The Three-Tier Progressive Disclosure of Skills (The Skill Iceberg). By partitioning domain knowledge into lightweight frontmatter metadata ($P = 1.0$), on-demand instruction manuals ($P \approx 10^{-1}$), and specialized execution scripts ($P \approx 10^{-3} - 10^{-4}$), agent harnesses achieve orders-of-magnitude reduction in context consumption while maintaining complete access to extensive tooling libraries.

4.3 Lifecycle Hooks and Supervisory Control

Harness lifecycle hooks intercept agent control flow at critical execution boundaries. Formally, lifecycle hooks operate as a **Ramadge–Wonham Supervisory Controller** (S/G).

Let the unconstrained agent execution be modeled as a discrete-event generator $G = \langle X, \Sigma, \delta, x_0, X_m \rangle$, where events $\Sigma = \Sigma_c \cup \Sigma_u$ are partitioned into controllable events Σ_c (tool invocations, file writes, network requests) and uncontrollable events Σ_u (model token generation, host hardware interrupts, external process termination). A supervisory hook $S : L(G) \rightarrow \mathcal{P}(\Sigma)$ dynamically disables tool invocations that violate safety invariants:

$$L(S/G) \subseteq K \subset L(G) \quad (22)$$

where K is the formal language of safe execution traces (e.g., preventing writes to protected branches, restricting filesystem mutations to leased workspaces, or halting when spending caps are exceeded).

Figure ?? formalizes the eight-stage agent execution pipeline and maps the concrete lifecycle hook implementations across leading frontier agent engines: Anthropic’s Claude Code, OpenAI’s Codex/Swarm, Google’s Gemini Agent Development Kit (ADK), and the Model Context Protocol (MCP).

4.3.1 What Lifecycle Hooks Allow and What They Exclude

Understanding the mathematical boundary of lifecycle hooks is crucial for operating kernel design:

- **What Hooks Allow:**




Figure 7: The Agent Lifecycle Hook pipeline and comparative interception points across Claude Code, OpenAI Codex, Google Gemini ADK, and Model Context Protocol (MCP). The Pre-Tool Interception Gate acts as a Ramadge–Wonham supervisory controller (S/G), providing the critical effect-boundary enforcement point capable of halting or rewriting unsafe tool calls before OS execution.

- (a) *Deterministic Pre-Execution Veto*: The `PreToolCall` hook can inspect parameters, evaluate capability macaroons, verify file leases, and abort execution before any physical syscall occurs.
- (b) *Parameter Normalization and Containment*: Hooks can rewrite paths to confine edits within a designated sandbox, enforcing confinement at the API boundary.
- (c) *Confidentiality Scrubbing*: The `PostToolCall` hook can intercept stdout/stderr, scrubbing discovered API tokens and private credentials before they enter the model’s KV-cache.
- (d) *Spend and Time Metering*: Hooks enforce hard stopping conditions on cumulative token usage and execution time.

• **What Hooks Cannot Allow (The Inherent Limitations):**

- (a) *Zero Control Over Internal Deliberation*: Hooks intercept only external events (Σ_c). They cannot steer the internal attention weights of the transformer during autoregressive generation.
- (b) *Cannot Repair Corrupted KV-Activations*: Once an untrusted string or misleading prompt injection enters the context buffer, hooks cannot selectively excise its semantic influence without rolling back the entire context window.
- (c) *Post-Hoc Ineffectiveness*: Intercepting an event after execution (`PostToolCall`) cannot undo destructive side effects (such as an unauthenticated network egress or an unbacked drop of a production table) that already occurred in the host environment.

5 Multi-Agent Swarm Topologies and Interaction Patterns

When multiple agents collaborate on a codebase, their interaction follows standard topological patterns:

Topology	Coordination Characteristics
Hierarchical Orchestrator	A primary agent decomposes goals, spawns specialized subagents, aggregates findings, and enforces review gates. Single point of failure; prone to orchestrator context saturation.
Peer-to-Peer Gossip	Agents communicate directly via broadcast or targeted channels. Highly resilient, but vulnerable to message explosion ($O(N^2)$) and consensus divergence.
Tuple Space (Blackboard)	Agents coordinate asynchronously through an associative shared memory (Linda model). Writes, reads, and takes are atomic; enables loose temporal and spatial coupling.
Market-Based Clearing	Tasks and resources are allocated through sealed-bid or English auctions under budget constraints. Prevents resource hoarding and aligns allocation with task priority.

Table 5: Taxonomy of multi-agent swarm interaction topologies.

Figure ?? contrasts these five communication patterns geometrically, highlighting the structural failure modes of peer chatter versus the resilience of kernel-mediated stigmergy.




Figure 8: Taxonomy of Multi-Agent Swarm Topologies. While peer-to-peer swarms suffer from $O(N^2)$ message inflation and consensus failure, the Stigmergic Blackboard coordinates agents via environmental state mutations mediated by the authoritative operating kernel.

5.1 Stigmergy and Shared Coordination Mediums

Why do multi-agent chat systems fail to scale past three or four concurrent participants? The answer lies in the physics of communication channels. In 1959, French biologist Pierre-Paul Grassé coined the term **Stigmergy** (from the Greek *stigma*, mark, and *ergon*, work) to explain how social insects construct intricate, cathedral-like termite mounds without central leadership or point-to-point communication.

A termite does not whisper instructions to its neighbors. Instead, it deposits a moist pellet of earth impregnated with a chemical pheromone. Nearby termites sense the physical deposit and the pheromone concentration gradient, which stimulates them to deposit their own pellets at that exact location. The physical modification of the environment regulates subsequent behavior: *the environment itself serves as the coordination medium*.

In naive multi-agent software systems, coordination is attempted through direct conversational messaging (chat swarms). For N autonomous agents, the number of potential pairwise

communication edges scales as:

$$E = \binom{N}{2} = \frac{N(N-1)}{2} = \mathcal{O}(N^2) \quad (23)$$

Under direct messaging, every agent’s context window becomes saturated with conversational chatter: acknowledgments, status queries, redundant progress reports, and semantic misinterpretations. Worse, because statistical language models possess non-zero hallucination rates ($\epsilon > 0$), unanchored peer chat causes errors to compound exponentially: agent *A* hallucinates an assumption; agent *B* accepts it and extrapolates; agent *C* builds upon *B*. Within three turns, the swarm descends into mutual conviction around a completely fabricated state—a pathology termed **Agentic Psychosis**.

In an Agent Operating Kernel, swarm coordination is strictly **Stigmergic**:

1. **The Git Worktree as the Physical Substrate:** Agents do not tell peers which lines of code they modified. They commit concrete diffs to isolated worktree branches. Other agents inspect the commit graph and the physical AST on disk.
2. **The Single-Writer WAL as the Event Horizon:** State mutations (task claims, test completions, lock leases) are written to a single-writer write-ahead log. Agents observe state transitions by streaming deterministic change events from the kernel, never by polling peers.
3. **Associative Tuple Spaces (Linda Model):** Agents coordinate asynchronously via atomic Linda tuples:

$$\langle \text{tag}, \text{key}, \text{payload} \rangle \quad (24)$$

An orchestrator posts $\langle \text{task_dispatch}, \text{refactor_auth}, \text{spec} \rangle$. An available worker takes the tuple atomically, computes the change, and writes $\langle \text{task_result}, \text{refactor_auth}, \text{diff_sha} \rangle$. Neither agent needs to know the other’s process identity, physical location, or memory state.

4. **Mutual Exclusion Locks and Port Registries:** Contention over scarce resources (TCP ports, shared database tables, benchmark runners) is resolved by atomic kernel mutexes. If agent *A* holds lock $\mathcal{L}_{\text{db}}$, agent *B*’s attempt to acquire it fails immediately at the kernel gate, preventing silent data corruption.

By replacing conversational chatter with physical substrate mutations, stigmergic coordination scales linearly $\mathcal{O}(N)$ with fleet size, eliminates split-brain hallucinations, and preserves deterministic replayability.

5.2 Context Compaction, Externalized Plans, and Cryptographic Identity

Scaling autonomous agents across multi-day engineering campaigns requires solving three inter-related state problems: context window limits, planning durability, and persistent identity.

5.2.1 The Failure of Naive Context Compaction

When an agent’s context window approaches capacity ($T \rightarrow T_{\text{max}}$), execution engines trigger context compaction. In unmediated harnesses, compaction is implemented as a lossy prompt summarization: the model is instructed to “summarize the conversation history and key achievements so far.”

For rigorous software engineering, naive prose summarization is catastrophic:

- **Precision Erasure:** Exact 40-character git commit SHAs, file path basenames, compiler error line numbers, and active lock identifiers are compressed into vague natural language summaries (e.g., “I refactored the auth module and encountered some test issues”).

- **Hallucinatory Success Promotion:** In subsequent turns, the model reads its own positive summary and hallucinates that failing tests passed or that uncommitted edits were merged.
- **Obligation Evaporation:** Active normative obligations, pending PR reviews, and unreleased capability leases vanish from context, leaving orphaned locks and unfinished transactions in the host system.

The operating kernel replaces lossy summarization with **Structured State-Delta Projection**:

$$\mathcal{P}_{\text{compact}}(h_t) = \langle \text{HEAD}_{\text{SHA}}, \mathcal{L}_{\text{active}}, \mathcal{C}_{\text{uncommitted}}, \mathcal{O}_{\text{pending}}, \mathcal{R}_{\text{verified}} \rangle \quad (25)$$

The kernel extracts the exact machine-checkable invariants: the currently checked-out commit hash, active lock leases with expiration timestamps, dirty file paths, open obligation receipts, and passing test hashes. This structured ledger is re-injected as a verified system preamble, guaranteeing zero loss of physical state across compaction boundaries.

5.2.2 Externalized Durable Plans as State Machines

In naive agent frameworks, a plan is an ephemeral Markdown checklist generated in the model’s scratchpad. As execution progresses and tokens scroll out of the attention window, the agent suffers from *intention drift*: it forgets earlier subgoals, repeats already-completed steps, or prematurely declares success.

In an accountable architecture, a plan is an **Externalized State Machine** managed by the kernel. The plan is stored in the database as a Directed Acyclic Graph (DAG) of discrete tasks with explicit dependency edges. Task transitions are executed via atomic Compare-And-Swap (CAS) operations:

$$\text{CAS}(\text{task_id}, \text{expected} = \text{IN_PROGRESS}, \text{new} = \text{COMPLETED}, \text{receipt} = \mathcal{H}_{\text{test}}) \quad (26)$$

An agent cannot unilaterally check off a plan item merely by writing prose; it must submit a verified cryptographic receipt (such as a passing test run digest or a signed git commit) to advance the state machine. If an agent process crashes or is killed mid-task, the externalized plan state remains intact in the kernel WAL, allowing a replacement worker to resume immediately from the exact point of interruption.

5.2.3 Persistent Identity Data Structures

Finally, we formalize the separation of durable identity from transient execution bodies:

- **Cryptographic Keypairs (Ed25519):** Each durable agent identity $\mathcal{I}$ is bound to an asymmetric keypair $(sk_{\mathcal{I}}, pk_{\mathcal{I}})$. All notes, commit assertions, and task proposals are signed with $sk_{\mathcal{I}}$.
- **Attenuated Capabilities (Macaroons):** Authorization is carried via chained capability tokens:

$$M_n = \mathcal{H}(M_{n-1} \parallel c_n) \quad (27)$$

where c_n represents a contextual caveat (e.g., `path_prefix = src/auth/`, `max_spend = $5.00`, `expires = 1774320000`). Caveats can only narrow authority, never widen it.

- **Tamper-Evident Event Chains:** Every action and observation is chained into an append-only Merkle tree:

$$R_t = \mathcal{H}(R_{t-1} \parallel \mathcal{A}_{\text{act}} \parallel \mathcal{H}(\mathcal{O}_{\text{obs}}) \parallel \text{timestamp}) \quad (28)$$

providing non-repudiable audit trails that survive process recycling, container termination, and machine reboots.

5.3 Durable Identity vs. Ephemeral Execution Bodies

A fundamental architectural principle of robust swarms is the strict decoupling of **Durable Identity** from **Ephemeral Execution Bodies**:

- An *Identity* represents the persistent persona, cryptographic keypair, credit balance, and accountability record of an agent. Identity survives machine reboots and process terminations.
- A *Body* is a temporary OS subprocess or container leased to execute a specific work slice.

When a subprocess crashes or hangs, the kernel reclaims the ephemeral body via **Zombie Salvage** (§6), returning unfinished task claims to the work register and assigning a fresh body to resume the durable identity’s mission.

6 The Agent Harness Lifecycle

An LLM is a stateless mathematical function mapping an input sequence of tokens to a probability distribution. An *Agent* is formed only when this stateless sampler is wrapped in a stateful execution loop known as the **Agent Harness**.

Figure 1 illustrates the closed agentic loop, stratified into three architectural tiers: the probabilistic Deliberation Layer, the confined Execution Layer, and the deterministic Governance Layer.




Figure 9: The closed agentic architecture loop: three-layer Swiss stratification separating probabilistic deliberation from deterministic kernel coordination.

The harness operates as a deterministic finite-state machine driving an indeterminate generator:

1. **Pre-Invocation Assembly:** The harness compiles system instructions, tool schemas, retrieved repository symbols, and recent conversation turns into the token buffer.
2. **Inference:** The prompt is transmitted to the inference engine. Tokens stream autoregressively until a function-call termination token is detected.

3. **Validation and Policy Gate:** The parser extracts tool arguments. The supervisor validates types against JSON-Schema definitions and verifies authority against local capability policies.
4. **Sandboxed Execution:** The command executes inside an isolated runtime container (macOS Seatbelt, Linux Bubblewrap) under strict egress caps.
5. **Kernel Witnessing:** Mutations, locks, and file reservations are submitted to the local kernel, appending a receipt to the audit log.
6. **Observation Injection:** Raw output is sanitized, truncated to preserve token budget, and appended to the context buffer for the next deliberation step.

7 Compute Topology: Frontier Cloud vs. Edge Actuation

Autonomous agent fleets are split across a fundamental physical asymmetry:

Attribute	Frontier Cloud Reasoning	Local Kernel / Host Edge
Workload	Multi-step decomposition, code synthesis	Deterministic ordering, file I/O, compilation
Latency	500 ms – 10 s per deliberation step	< 1 ms (POSIX locks, SQLite WAL writes)
Cost	\$1.00 – \$15.00 per million tokens	Local CPU cycles (effectively zero marginal cost)
Trust Boundary	Third-party cloud datacenter	Operator-controlled workstation / secure enclave
Failure Mode	Hallucination, semantic drift, timeout	Crash-stop, lock contention, disk exhaustion

Table 6: Asymmetry between cloud inference engines and local host execution surfaces.

Because cloud reasoning is high-latency and probabilistic, it is physically incapable of resolving local race conditions. If two cloud agents concurrently attempt to edit line 42 of the same file, relying on their internal world models to avoid collisions guarantees data corruption. The resolution must be executed locally by a synchronous, zero-latency authority.

8 The Kernel Gap: Why LLMs Cannot Be Their Own OS

A seductive but fatal fallacy in modern AI engineering is the belief that an LLM can serve as its own operating system kernel—that prompts such as “*Please do not modify files currently being edited by other agents*” are sufficient to preserve system integrity.

We term this discrepancy the **Kernel Gap**. Language models fail as operating systems across four fundamental dimensions:

8.1 1. Non-Determinism vs. Invariant Regimentation

An operating system kernel must enforce *regimented invariants*: physical guarantees that cannot be bypassed by any user process (e.g., memory page protection via MMU hardware). Prompts are mere suggestions. Even fine-tuned models exhibit a non-zero probability of instruction failure, hallucination, or adversarial prompt injection. A security policy enforced solely by a prompt is an advisory boundary, not an invariant.

8.2 2. Concurrency and the Mutual Exclusion Anomaly

In any concurrent system with two or more uncoordinated writers, atomic mutual exclusion is required to prevent race conditions. Without an external serialized lock broker, concurrent LLM agents exhibit the classic lost-update anomaly: Agent A reads version v_0 , Agent B reads v_0 , Agent A writes v_A , and Agent B subsequently overwrites with v_B , silently destroying Agent A’s work without compiler or runtime warning.

8.3 3. Attenuation and Privilege Escalation

In classical capability-based operating systems (e.g., KeyKOS, EROS, Capsicum), authority flows monotonically downward: a child process can never possess more capability than its parent. When an LLM delegates sub-tasks by synthesizing new prompt instructions, authority frequently inflates due to broad, imprecise system instructions. True privilege containment requires cryptographically unforgeable, attenuated tokens (Harbor Cards, Chapter 2).

8.4 4. Mortality and Economic Accountability

Subprocesses are transient. An agent can crash, get SIGKILLed, or reset its context, discarding all local history. If reputation, debt, and liability exist only in process RAM, defecting agents face zero consequence for bad actions. Accountability requires persistent outcome ledgers, Merkle-backed event histories, and economic escrows that survive the death of the acting process (Chapters 5–7).

9 Chapter Map and Structure of the Book

The remainder of this textbook is organized into eight cumulative chapters across four distinct operational levels:

- **Chapter 1: The Single-Writer Kernel (Ground Truth).** We construct the minimal deterministic local kernel using SQLite WAL mode, establishing serializability, synchronous state decidability, and crash durability across process failures.
- **Chapter 2: The Anchor Protocol (Capability Delegation).** We design cryptographically verifiable capability tokens (Harbor Cards) and prove that delegated authority can only narrow at every hop.
- **Chapter 3: The Sealed Harbor (Confinement & Leakage).** We construct isolated execution harbors that bound information leakage between distrustful principals under strict egress quotas.
- **Chapter 4: The Legible Swarm (Human Supervision).** We derive the information-theoretic cost of supervisory attention, proving how digests must retain evidence paths without drowning the operator in raw telemetry.
- **Chapter 5: From Spawn to Person (Durable Agency).** We define persistent agent identity, append-only outcome ledgers, and reputation mechanisms that survive process restarts.
- **Chapter 6: The Harbor Economy (Markets & Clearing).** We establish bilateral clearing, underwriting reserves, and dynamic pricing for scarce host resources.
- **Chapter 7: The Bonded Commons (Game-Theoretic Governance).** We formalize claim signaling as an incentive-compatible repeated game, proving when cooperation is sustained under bonded escrow.

- **Chapter 8: The Federated Harbor (Cross-Machine Consensus).** We extend governance across distrustful host boundaries via witness logs and sheaf-theoretic consistency without central sovereignty.

References

- [1] G. Agha. *Actors: A Model of Concurrent Computation in Distributed Systems*. MIT Press, Cambridge, MA, 1986. 2
- [2] J. L. Austin. *How to Do Things with Words*. Oxford University Press, 1962. 2
- [3] R. H. Bordini, J. F. Hübner, and M. Wooldridge. *Programming Multi-Agent Systems in AgentSpeak using Jason*. John Wiley & Sons, 2007.
- [4] M. E. Bratman. *Intention, Plans, and Practical Reason*. Harvard University Press, Cambridge, MA, 1987. 2
- [5] T. Charrier, F. Ouchet, and F. Schwarzentruher. Big Brother logic: Reasoning about agents equipped with surveillance cameras in the plane. In *Proc. 13th International Conference on Autonomous Agents and Multiagent Systems (AAMAS)*, pages 421–428, 2014.
- [6] Foundation for Intelligent Physical Agents. FIPA ACL Communicative Act Library Specification. Technical Report SC00037J, FIPA, 2002. 2
- [7] C. Hewitt, P. Bishop, and R. Steiger. A universal modular ACTOR formalism for artificial intelligence. In *Proc. 3rd International Joint Conference on Artificial Intelligence (IJCAI)*, pages 235–245, 1973. 2
- [8] C. A. R. Hoare. Communicating sequential processes. *Communications of the ACM*, 21(8):666–677, 1978.
- [9] N. F. Liu, K. Lin, J. Hewitt, A. Paranjape, M. Bevilacqua, F. Petroni, and P. Liang. Lost in the middle: How language models use long contexts. *Transactions of the Association for Computational Linguistics*, 12:157–173, 2024. 3
- [10] A. S. Rao and M. P. Georgeff. Modeling rational agents within a BDI-architecture. In *Proc. 2nd International Conference on Principles of Knowledge Representation and Reasoning (KR)*, pages 473–484, 1991. 2
- [11] A. S. Rao and M. P. Georgeff. BDI agents: From theory to practice. In *Proc. 1st International Conference on Multi-Agent Systems (ICMAS)*, pages 312–319, 1995. 2
- [12] J. R. Searle. *Speech Acts: An Essay in the Philosophy of Language*. Cambridge University Press, 1969. 2
- [13] N. Shinn, F. Cassin, B. Labash, A. Gopinath, P. Narasimhan, and S. Yao. Reflexion: Language agents with verbal reinforcement learning. In *Proc. NeurIPS*, 2023.
- [14] Y. Shoham and K. Leyton-Brown. *Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations*. Cambridge University Press, 2009.
- [15] R. G. Smith. The Contract Net Protocol: High-level communication and control in a distributed problem solver. *IEEE Transactions on Computers*, C-29(12):1104–1113, 1980.
- [16] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting in language models. In *Proc. ICLR*, 2023.

- [17] S. Yao, D. Yu, J. Zhao, I. Shafran, T. L. Griffiths, Y. Cao, and K. Narasimhan. Tree of Thoughts: Deliberate problem solving with large language models. In *Proc. NeurIPS*, 2023.